

1 High-level, high-resolution ocean modeling at all scales

2 **Gregory L. Wagner¹, Simone Silvestri¹, Navid C. Constantinou^{2,3},**
 3 **Jean-Michel Campin¹, Chris Hill¹, Ali Ramadhan⁴, Tomás Chor⁵,**
 4 **Jago Strong-Wright⁶, Xin Kai Lee¹, Francis Poulin⁷, Keaton J. Burns¹,**
 5 **John Marshall¹, and Raffaele Ferrari¹**

6 ¹Massachusetts Institute of Technology, Cambridge, MA, USA

7 ²University of Melbourne, Parkville, VIC, Australia

8 ³ARC Center of Excellence for the Weather of the 21st Century, Australia

9 ⁴atdepth MRV Inc., 82 Wendell Ave, STE 100, Pittsfield, MA 01201, USA

10 ⁵Department of Atmospheric and Oceanic Science, University of Maryland, College Park, MD, USA

11 ⁶DAMTP, University of Cambridge, United Kingdom

12 ⁷University of Waterloo, Waterloo, ON, Canada

13 Abstract

14 We describe the user interface, governing equations, and numerical methods that underpin
 15 new ocean modeling software called “Oceananigans”. Oceananigans’ design realizes a vision
 16 for Earth system modeling that combines three synergistic elements: *(i)* a relatively simple
 17 finite volume algorithm *(ii)* optimized for high-resolution simulations on GPUs and *(iii)*
 18 implemented in a high-level programming language (the Julia programming language in
 19 our case) behind a productive user interface. The outcome is a modeling system primed
 20 for rapid community development that *also* pursues state-of-the-art fidelity. Capability
 21 comparable to more complex operational software is achieved via novel numerical methods
 22 and “brute force” high resolution. We argue that our already-significant achievements in
 23 parameterization, numerical methods, and model development demonstrate the potential of
 24 this vision for accelerating progress in computational Earth science.

25 Plain Language Summary

26 We describe the user interface, governing equations, and numerical methods that underpin
 27 new ocean modeling software called “Oceananigans”. Oceananigans is written in
 28 the high-level Julia programming language, making it much easier to install, learn, and use
 29 than existing modeling software. Unlike most ocean modeling software, Oceananigans can
 30 be used to simulate motions over a wide range of spatial scales depending on the problem
 31 being considered — from the long durations and planetary scales useful for weather and
 32 climate modeling, to sub-second, sub-meter scales useful for high-fidelity simulations in lim-
 33 ited domains. In addition to being flexible and ease-to-use, Oceananigans is *also* the fastest
 34 ocean modeling software today, because it utilizes GPUs. Each of these three achievements
 35 is noteworthy on their own. But taken together, we argue that these achievements indicate
 36 the potential our modeling strategy has to accelerate progress in computational Earth
 37 science .

38 1 Introduction

39 Computation is fundamental to ocean and climate science, such that software is rate-
 40 limiting for scientific progress. Since the first general circulation models ran on primitive
 41 computers (Phillips, 1956; Bryan, 1969), advances in hardware, numerical methods, and the
 42 approximate parameterization of otherwise unresolved processes have improved the fidelity
 43 of ocean simulations (Griffies et al., 2015). Yet as technology advances, the gap between
 44 potential and practice in ocean modeling is stagnant or widening, to the point that most
 45 software today *(i)* can no longer use the world’s fastest computers, *(ii)* relies on outdated,
 46 inefficient user interfaces, and *(iii)* is still useful for only a limited subset of the wide variety
 47 of ocean modeling problems.

This paper describes new ocean modeling software written in the Julia programming language (Bezanson et al., 2017). Our software features an innovative user interface, can simulate nonhydrostatic and hydrostatic motion at any scale using novel numerical methods and parameterizations, and runs on laptops and clusters of GPUs. Through effective use of GPUs, our software runs up to 10–50× faster than CPU-based software using equivalent resources. The software is being developed by the Climate Modeling Alliance as part of a larger effort to develop a climate model automatically-calibrated to observations and high resolution simulations, and with quantified uncertainty. The software is called “Oceananigans”.

Our work shows the potential of a strategy that pairs simple algorithms written in high-level languages with GPU-enabled high-resolution capability to achieve accuracy with accessible code. Our goal is to accelerate the *process* of numerical method and parameterization development — and through a longer process of collective effort, accelerate progress in ocean and climate science.

1.1 From millimeters to millenia

The evolution of ocean circulation over millennia is controlled by turbulent mixing with scales that range down to millimeters. Two distinct systems have evolved to model and understand this huge range of oceanic motion: “GCMs” (general circulation models) for hydrostatic regional-to-global scale simulations, and simpler software for nonhydrostatic, meter-scale large eddy simulations (LESs) that are high-fidelity but limited in duration and extent. Compared to LES, GCMs usually invoke more elaborate numerical methods and parameterizations to cope with the global ocean’s complex geometry and the more significant impacts of unresolved subgrid processes.

Oceananigans began as software for LES on GPUs (Ramadhan et al., 2020), by perfecting an approach for hybrid hydrostatic/nonhydrostatic dynamical cores pioneered by MITgcm (Marshall, Adcroft, et al., 1997). Our simple nonhydrostatic LES algorithm was then adapted and optimized for a hydrostatic GCM on GPUs, achieving unprecedented computational performance (Silvestri, Wagner, Constantinou, et al., 2024). At the same time, we developed high-quality, LES-inspired numerical methods for turbulence-resolving simulations (Silvestri, Wagner, Campin, et al., 2024) — resulting in a system suited to a brute force, resolution-focused approach to accurate simulations. This “LES the ocean” strategy is appealingly simple compared to alternatives relying on generalized vertical coordinates (Shchepetkin & McWilliams, 2005; Leclair & Madec, 2011; Petersen et al., 2015), Lagrangian vertical advection (Halliwell, 2004; Griffies et al., 2020), or unstructured horizontal grids (Ringler et al., 2013; Danilov et al., 2017; Korn et al., 2022). We hypothesize that “resolution everywhere” alleviates the need for more targeted resolution through unstructured grids and will reduce the spurious numerical mixing that pollutes the fidelity of lower-resolution simulations (Griffies et al., 2000), while yielding a plethora of additional improvements (Chassignet & Xu, 2017, 2021; Kiss et al., 2020). At the same time, using simple algorithms preserves the accessibility of our source code and maximizes the benefits of the Julia programming language.

1.2 Why programmable interfaces matter

In 1984, Cox published the first description of generalizable ocean modeling software (Cox, 1984; Griffies et al., 2015). The “Cox model” is written¹ in FORTRAN 77 and features a multi-step user interface for building new models: first, source code modifications are written to determine, for example, domain geometry and boundary conditions, emplaced into the “base code”, and compiled. Next, a text-based namelist file is used to determine

¹ The Cox Model is also designed for a single machine, the Control Data Corporations’ Cyber 205.

95 parameters like the stop iteration, mixing coefficients, and solver convergence criteria. Cox
 96 (1984) provided three example model configurations to illustrate the user interface.

97 With forty years of progress in software engineering, numerical methods, and parameterization
 98 of unresolved processes, and more than a billion times² more computational power, today's ocean models bear little resemblance to the Cox model — *except* for their
 99 user interfaces. Current interfaces, though obviously more advanced than Cox's, still impose
 100 multi-step workflows that invoke several programming paradigms. These multi-step
 101 workflows typically require the generation of input data using a separate scripting language,
 102 configuration of numerous namelists, and source code modifications to change the model
 103 equations in ways not accessible through a change of parameters.
 104

105 Multi-step workflows are inefficient. One of our most important contributions is the
 106 development of a fundamentally different, programmable user interface that provides a
 107 seamless workflow for numerical experiments including setup, execution, analysis, and
 108 visualization using a single script. Programmable interfaces written in scripting languages
 109 like Python and Julia are the interface of choice and engine of progress in countless fields
 110 from visualization to machine learning, and their benefits transfer to ocean modeling. A
 111 particularly inspiring example of a productive user interface for computational physics is
 112 provided by the CPU-based, spectral Dedalus (Burns et al., 2020) framework for solving
 113 partial differential equations in simple geometries.

114 A programmable interface shines for simple problems — but doesn't just help new
 115 users. More importantly, this workflow accelerates the implementation of new numerical
 116 methods and parameterizations by experienced developers. It facilitates writing and
 117 relentlessly refactoring comprehensive test suites. It enables fast prototyping with tight
 118 implementation-evaluation iterations. It makes it easier to collaborate by communicating
 119 concise but evocative code snippets. It makes Oceananigans fun to use. Leveraging this
 120 programmable interface together with the intrinsic productivity of the Julia programming
 121 language, Oceananigans has progressed from a simple system for serial nonhydrostatic
 122 modeling (Ramadhan et al., 2020) to parallelized software with capabilities at all scales up
 123 to global hydrostatic simulations with breakthrough performance (Silvestri, Wagner,
 124 Constantinou, et al., 2024), using innovative numerical methods (Silvestri, Wagner, Campin,
 125 et al., 2024) and new, automatically-calibrated vertical mixing parameterizations (Wagner et
 126 al., 2024). Users benefit too.

127 The Julia programming language — a programming language suitable for both scripting
 128 and applications that is at once compiled and fast, but also dynamic and productive — has
 129 a lot to do with the feasibility of our design. Unlike functions in pure Python, for example,
 130 Julia functions implemented by users for forcing and boundary conditions can operate even
 131 in high performance contexts on GPUs. Julia enables unique Oceananigans features, such as
 132 interactivity, extensibility, automatic installation on any system, and portability to laptops
 133 and GPUs through advanced Julia community tools (Besard et al., 2018; Churavy, 2024).
 134 Oceananigans' achieves breakthrough performance by using GPUs, but remains accessible
 135 to students using Windows laptops. Easy installation on personal computers facilitates
 136 creative computation, since complex numerical experiments can be prototyped productively
 137 in a comfortable personal environment before transferred to a high performance environment
 138 for production runs.

139 Productive interfaces are only as powerful as the capability they expose. Oceananigans'
 140 capabilities combine a range of capabilities offered by other systems: a design for modeling
 141 across scales from MITgcm (Marshall, Adcroft, et al., 1997; Marshall, Hill, et al., 1997),

² Compare the Cyber 205's 1 GFlop (https://www.computinghistory.org.uk/userdata/files/the-control-data-cyber-205-computer-system-the_practical-supercomputer.pdf) to Aurora's 1 exaflop ([https://en.wikipedia.org/wiki/Aurora_\(supercomputer\)](https://en.wikipedia.org/wiki/Aurora_(supercomputer)))

142 a simple and performant algorithm for LES from PALM and pycles (Pressel et al., 2015),
 143 GPU capabilities like Veros (Häfner et al., 2021), and a high-level interface like Thetis
 144 (Kärnä et al., 2018). Oceananigans assembles these diverse features behind a productive
 145 programmable interface.

146 1.3 Outline of this paper

147 This paper provides a tutorial-style introduction to Oceananigans' user interface, gov-
 148 erning equations, and finite volume discretization. Example simulations that illustrate
 149 Oceananigans' capabilities using visualizations, paired with code snippets that generate
 150 them, are interspersed through the paper. Our aim is to evidence and explain Oceananig-
 151 gans tripartite achievement: performance, flexibility, and friendliness at the same time.
 152 For a more detailed and comprehensive reference to the governing equations and numer-
 153 ical methods used by Oceananigans we refer the reader to our online documentation at
<https://clima.github.io/OceananigansDocumentation>.

155 Section 2 begins by explicating the basic innovations of Oceananigans' programmable
 156 interface section 2 using two example simulations: two-dimensional turbulence, and a forced
 157 passive tracer advected by two-dimensional turbulence. In section 3, we write down the
 158 governing equations that underpin Oceananigans' nonhydrostatic and hydrostatic models.
 159 We build up our presentation with examples that progress from simple direct numerical
 160 simulations of freshwater cabbeling and flow around a cylinder, to a 1/12th degree eddying
 161 global ocean simulation.

162 In section 4, we provide a primer to the finite volume spatial discretization that
 163 Oceananigans uses to solve the nonhydrostatic and hydrostatic equations. This section estab-
 164 lishes Oceananigans' unique suitability for turbulence-resolving simulations that have mini-
 165 mal, implicitly dissipative advection schemes based on Weighted Essentially Non-Oscillatory
 166 (WENO) reconstruction. We conclude in section 6 by outlining future development work
 167 and anticipating the next major innovations in ocean modeling which, we hope, will someday
 168 render the present work obsolete. Appendix A describes the time-stepping scheme.

169 2 Oceananigans, the library

170 Oceananigans is fundamentally a *library* of tools for building models by writing pro-
 171 grams called “scripts”. This departs from the usual framework wherein software provides
 172 pre-written monolithic programs that are configured with parameters. For writing scripts,
 173 Oceananigans provides syntax that judiciously combines mathematical symbols with nat-
 174 ural language, including acronyms when appropriate. Our design enables clear, evocative
 175 scripting that approaches the effectiveness of writing in communicating computational sci-
 176 ence.

177 2.1 Hello, ocean

178 The only way to learn new ocean modeling software is by building simulations with it.
 179 Our first example in listing 1 sets up, runs, and visualizes a simulation of two-dimensional
 180 turbulence. The 22 lines of listing 1 illustrate one of Oceananigans' main achievements: a
 181 numerical experiment may be completely described by a single script. To execute the code
 182 in listing 1 it is copied into a file (call this `hello_ocean.jl`) and executed by typing `julia`
 183 `hello_ocean.jl` at a terminal.

```
184
185 1 using Oceananigans
186 2
187 3 # The third dimension is "flattened" to reduce the domain from three to two dimensions.
188 4 topology = (Periodic, Periodic, Flat)
189 5 architecture = GPU() # CPU() works just fine too for this small example.
190 6 x = y = (0, 2π)
191 7 grid = RectilinearGrid(architecture, size=(256, 256), x, y, topology)
```

```

193
194     8
195     9 model = NonhydrostaticModel(; grid, advection=WENO(order=9))
196    10
197    11 ε(x, y) = 2rand() - 1 # Uniformly-distributed random numbers between [-1, 1].
198    12 set!(model, u=ε, v=ε)
199    13
200    14 simulation = Simulation(model; Δt=0.01, stop_time=10)
201    15 run!(simulation)
202    16
203    17 u, v, w = model.velocities
204    18 ζ = Field(∂x(v) - ∂y(u))
205    19 compute!(ζ)
206    20
207    21 using CairoMakie
208    22 heatmap(ζ, colormap=:balance, axis=(; aspect=1))
209

```

Listing 1: A Julia script that uses Oceananigans and the Julia plotting library CairoMakie to set up, run, and visualize a simulation of two-dimensional turbulence on a Graphics Processing Unit (GPU). The initial velocity field, defined on line 9, consists of random numbers uniformly-distributed between -1 and 1 . The vorticity $\zeta = \partial_x v - \partial_y u$ is defined on line 16. The solution is visualized in figure 1.

Oceananigans scripts organize into four sections. The first three define the “grid” “model”, and “simulation”, and conclude with execution of the simulation. The fourth section, often implemented separately for complex or expensive simulations, performs post-processing and analysis. In listing 1, the grid defined on lines 4–7 determines the problem geometry, spatial resolution, and machine architecture. To use a CPU instead of a GPU, one writes `CPU()` in place of `GPU()` on line 5: no other changes to the script are required.

Lines 9–12 define the model, which solves the Navier–Stokes equations in two dimensions with a 9th-order Weighted, Essentially Non-Oscillatory (WENO) advection scheme (see section 4 for more information about WENO). The velocity components u, v are initialized with uniformly distributed random numbers between $[-1, 1]$. The model definition can also encompass forcing, boundary conditions, and the specification of additional terms in the momentum and tracer equations such as Coriolis forces or turbulence closures.

Line 14 builds a `Simulation` with a time-step $\Delta t = 0.01$ which will run until $t = 10$ (Oceananigans does not assume dimensionality by default, so time is non-dimensional via user input in this case). Later examples will illustrate how `Simulation` can be used to inject arbitrary user code into the time-stepping loop in order to log simulation progress or write output to disk. Lines 17–19 analyze the final state of the simulation by computing vorticity, illustrating Oceananigans’ toolbox for building expression trees of discrete calculus and arithmetic operations. The same tools may be used to define online diagnostics to be periodically computed and saved to disk while the simulation runs. Line 22 concludes the numerical experiment with a visualization. The result is shown in figure 1.

2.2 Integrating user code

With a programmable interface and aided by Julia’s just-in-time compilation, user functions specifying domain geometry, forcing, boundary conditions, and initial conditions can be incorporated directly into models without a separate programming environment. To illustrate function-based forcing, we modify listing 1 with code that adds a passive tracer which is forced by a moving source that depends on x, y, t . A visualization of the vorticity and tracer field generated by listings 1 and 2 are shown in figure 1.

```

238
239     1 @inline function circling_source(x, y, t)
240     2     δ, ω, r = 0.1, 2π/3, 2
241     3     dx = x + r * cos(ω * t)
242     4     dy = y + r * sin(ω * t)
243     5     return exp(-(dx^2 + dy^2) / 2δ^2)
244     6 end
245     7
246     8 forcing = NamedTuple(c => circling_source)
247     9 model = NonhydrostaticModel(; grid, advection=WENO(order=9), tracers=:c, forcing)
248
249

```

250

Listing 2: Implementation of a moving source of passive tracer with a function in a two-dimensional turbulence simulation. These lines of code replace the model definition on line 9 in listing 1. The prefix `@inline` helps to ensure `circling_source` successfully compiles on the GPU.

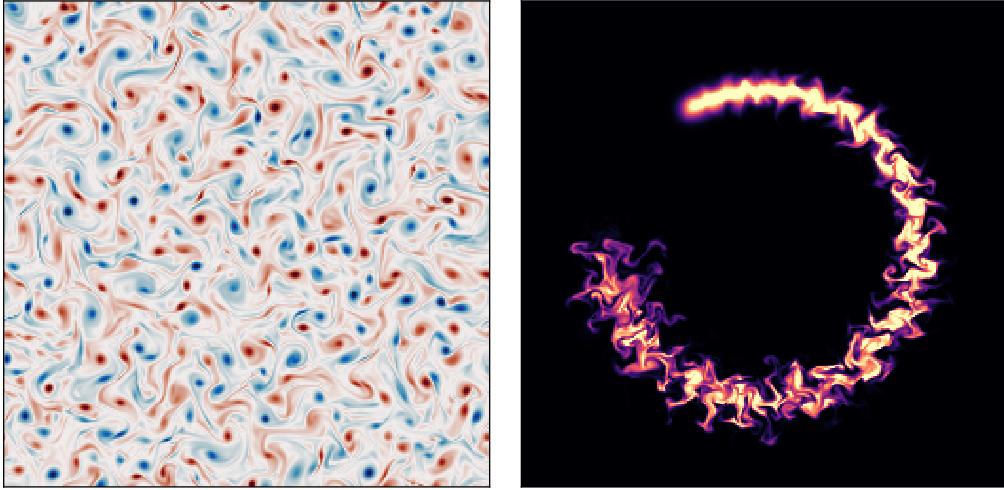


Figure 1: Vorticity after $t = 10$ (left) and a passive tracer injected by a moving source at $t = 2.5$ (right) in a simulation of two-dimensional turbulence using an implicitly-dissipative advection scheme.

251
252
253
254

Users can also insert arbitrary functions for more general tasks into the time-stepping loop. This supports things as mundane as printing a summary of the current model status or writing output, to more exotic tasks like nudging state variables or updating a diffusion coefficient based on an externally-implemented model.

255

2.3 Abstractions for arithmetic and discrete calculus

256
257
258
259
260
261
262

Abstractions representing unary, binary, and calculus operators produce a system for building lazy expression trees to be saved to disk during a simulation. Example calculations representing vorticity, $\zeta = \partial_x v - \partial_y u$, speed $s = \sqrt{u^2 + v^2}$, and the x -integral of enstrophy $Z = \int_0^{2\pi} \zeta^2 dx$ are shown in listing 3. Listing 3 illustrates the difference between lazy “operations” and concrete Fields that have “data” in memory for storing the result of a computation. Distinguishing between lazy operations and concrete fields avoids unnecessary memory allocation when implementing diagnostics.

263
264
265
266
267
268
269
270
271
272

```

1 u, v, w = model.velocities
2
3 # Lazy expression trees and reductions representing computations:
4 ζ = ∂x(v) - ∂y(u)
5 s = sqrt(u^2 + v^2)
6 Z = Integral(ζ^2, dims=1)

```

Listing 3: “Lazy” abstractions for expression trees and reductions — abstractions that *represent* computations to be performed at some future time as needed — support custom online diagnostics.

273 3 Governing equations and physical parameterizations

274 Oceananigans implements two “models” for ocean-flavored fluid dynamics: the HydrostaticFreeSurfaceModel, and the NonhydrostaticModel. Each represents a template for
 275 equations that govern the evolution of momentum and tracers. Both models are incompressible and make the Boussinesq approximation, which means that the density of the modeled
 276 fluid is decomposed into a constant reference ρ_0 and a small dynamic perturbation ρ' ,
 277

$$\rho(\mathbf{x}, t) = \rho_0 + \rho'(\mathbf{x}, t) \quad \text{where} \quad \rho' \ll \rho_0, \quad (1)$$

279 and $\mathbf{x} = (x, y, z)$ is position and t is time.

280 The relative smallness of ρ' reduces conservation of mass to a statement of incompressibility called the continuity equation,
 281

$$\nabla \cdot \mathbf{u} = 0, \quad (2)$$

282 where

$$\mathbf{u} \stackrel{\text{def}}{=} u \hat{\mathbf{x}} + v \hat{\mathbf{y}} + w \hat{\mathbf{z}}, \quad (3)$$

283 is the three-dimensional velocity field. Within the Boussinesq approximation, the momentum
 284 $\rho_0 \mathbf{u}$ varies only with the velocity \mathbf{u} . The effect of density variations is encapsulated by
 285 a buoyant acceleration,

$$b \stackrel{\text{def}}{=} -\frac{g\rho'}{\rho_0}, \quad (4)$$

286 where g is gravitational acceleration. The “buoyancy” b acts in the direction of gravity.

287 3.1 The NonhydrostaticModel

288 The NonhydrostaticModel represents the Boussinesq equations formulated *without* making
 289 the hydrostatic approximation typical to general circulation models. The Nonhydro-
 290 staticModel has a three-dimensional prognostic velocity field. Dynamic pressure P in the
 291 NonhydrostaticModel is written,

$$P = \rho_0 g z + \rho_0 p(\mathbf{x}, t), \quad (5)$$

292 where $\rho_0 g z$ is the static contribution to pressure and $\rho_0 p$ is the dynamic anomaly. p is called
 293 the kinematic pressure.

294 3.1.1 The NonhydrostaticModel momentum equation

295 The NonhydrostaticModel’s momentum equation incorporates advection by a back-
 296 ground velocity field, Coriolis forces, surface wave effects via the Craik-Leibovich asymptotic
 297 model (Craik & Leibovich, 1976; Huang, 1979), a buoyancy term allowed to be a nonlinear
 298 function of tracers and depth, a stress divergence derived from molecular friction or a tur-
 299 bulence closure, and a user-defined forcing term. The generic form of NonhydrostaticModel’s
 300 momentum equation is

$$\begin{aligned} \partial_t \mathbf{u} = & -\nabla p - \underbrace{(\mathbf{u} \cdot \nabla) \mathbf{u} - (\mathbf{u}_g \cdot \nabla) \mathbf{u} - (\mathbf{u} \cdot \nabla) \mathbf{u}_g}_{\text{advection}} - \underbrace{\mathbf{f} \times \mathbf{u}}_{\text{Coriolis}} \\ & + \underbrace{(\nabla \times \mathbf{u}_s) \times \mathbf{u} + \partial_t \mathbf{u}_s}_{\text{Stokes drift}} - \underbrace{b \hat{\mathbf{g}}}_{\text{buoyancy}} - \underbrace{\nabla \cdot \boldsymbol{\tau}}_{\text{closure}} + \underbrace{\mathbf{F}_u}_{\text{forcing}}, \end{aligned} \quad (6)$$

301 where \mathbf{u}_g is a prescribed “background” velocity field, p is the kinematic pressure, \mathbf{f} is the
 302 background vorticity associated with a rotating frame of reference, \mathbf{u}_s is the Stokes drift
 303 profile associated with a prescribed surface wave field, b is buoyancy, $\hat{\mathbf{g}}$ is the gravitational
 304 unit vector (usually pointing downwards, so $\hat{\mathbf{g}} = -\hat{\mathbf{z}}$), $\boldsymbol{\tau}$ is the stress tensor associated with
 305 molecular viscous or subgrid turbulent momentum transport, and \mathbf{F}_u is a body force.

To integrate equation (6) while enforcing (2), we use a pressure correction method that requires solving a three-dimensional Poisson equation to find p , which can be derived from $\nabla \cdot (6)$. This Poisson equation is often a computational bottleneck in curvilinear or irregular domains, and its elimination is the main motivation for making the hydrostatic approximation when formulating the HydrostaticFreeSurfaceModel, as described in section 3.2. We solve the Poisson equation using a fast, direct, FFT-based method (Schumann & Sweet, 1988), providing substantial acceleration over MITgcm’s conjugate gradient pressure solver (Marshall, Adcroft, et al., 1997). In irregular domains, we use a conjugate gradient iteration similar to MITgcm, except that we leverage the FFT-based solver as a preconditioner, which typically converges in fewer than 10 iterations. The pressure correction scheme is described further in appendix A2.

Using (2), advection in the NonhydrostaticModel is formulated in the “flux form”, which is conveniently expressed with indicial notation,

$$\text{advection} = u_j \partial_j u_i + u_{gj} \partial_j u_i + u_j \partial_j u_{gi} = \partial_j [(u_j + u_{gj}) u_i + u_j u_{gi}], \quad (7)$$

where, for example, the i -th component of the advection term is $[(\mathbf{u} \cdot \nabla) \mathbf{u}]_i = u_j \partial_j u_i$.

The formulation of the Stokes drift terms means that \mathbf{u} is the Lagrangian-mean velocity when Stokes drift effects are included (see, for example, Wagner et al., 2021). With a Lagrangian-mean formulation, equations (2) and (6) are consistent only when \mathbf{u}_s is non-divergent — or equivalently, when \mathbf{u}_s is obtained by projecting the divergence out of the usual Stokes drift (Vanneste & Young, 2022). As discussed by Wagner et al. (2021), the Lagrangian-mean formulation of (6) means that closures for LES strictly destroy kinetic energy, avoiding the inconsistency between resolved and subgrid fluxes affecting typical LES formulated in terms of the Eulerian-mean velocity (see also Pearson, 2018).

The labeled terms in (6) are controlled by arguments to NonhydrostaticModel invoked in both of listings 1 and 2. For example, “advection” chooses a numerical scheme to approximate the advection term in (6) and (7). As another example, we consider configuring the closure term in (6) to represent (i) molecular diffusion by a constant-coefficient Laplacian ScalarDiffusivity, (ii) turbulent stresses approximated by the SmagorinskyLilly eddy viscosity model (Smagorinsky, 1963; Lilly, 1983) for large eddy simulation, or (iii) omitting it entirely, which we use with WENO advection schemes (and which is also our default setting). In these three cases, the closure flux divergence $\nabla \cdot \boldsymbol{\tau} = \partial_m \tau_{nm}$ in indicial notation becomes

$$-\partial_m \tau_{nm} = \begin{cases} \partial_m (\nu \partial_m u_n) & \text{(ScalarDiffusivity)} \\ 0 & \text{(nothing)} \\ \partial_m \left(2 \underbrace{C_s \Delta^2 |\Sigma|}_{\nu_e} \Sigma_{nm} \right) & \text{(SmagorinskyLilly)} \end{cases} \quad (8)$$

where ν is the Laplacian diffusion coefficient, $\Sigma_{nm} = \partial_m u_n + \partial_n u_m$ is the strain tensor, $|\Sigma|$ is the magnitude of the strain rate tensor, C_s is the SmagorinskyLilly model constant, Δ scales with the local grid spacing, and ν_e is the eddy viscosity. (ScalarDiffusivity diffusion coefficients may also vary in time- and space. Other closure options include fourth-order ScalarBiharmonicDiffusivity, various flavors of DynamicSmagorinsky (Bou-Zeid et al., 2005), and the AnisotropicMinimumDissipation turbulence closure (Rozema et al., 2015; Vreugdenhil & Taylor, 2018) for large eddy simulations.)

Listing 4 implements a direct numerical simulation of uniform flow past a cylinder with no-slip boundary conditions, a molecular ScalarDiffusivity, and a centered second-order advection scheme. Lines 8–9 embed a cylindrical mask in a RectilinearGrid using a GridFittedBoundary, which generalizes to arbitrary three-dimensional shapes. The no-slip condition is implemented with ValueBoundaryCondition (a synonym for “Dirichlet” boundary conditions) on lines 14–15. Other choices include GradientBoundaryCondition (Neumann),

```

1 r, U, Re, Ny = 1/2, 1, 1000, 2048
2
3 grid = RectilinearGrid(GPU(), size=(2Ny, Ny), x=(-3, 21), y=(-6, 6),
4 topology=(Periodic, Bounded, Flat))
5
6 cylinder(x, y) = (x^2 + y^2) ≤ r^2
7 grid = ImmersedBoundaryGrid(grid, GridFittedBoundary(cylinder))
8
9 closure = ScalarDiffusivity(v=1/Re) # Try SmagorinskyLilly() for LES
10
11 no_slip = FieldBoundaryConditions(immersed=ValueBoundaryCondition(0))
12 boundary_conditions = (u=no_slip, v=no_slip)
13
14 # Implement a sponge layer on the right side of the domain that
15 # relaxes v → 0 and u → U over a region of thickness δ
16 @inline mask(x, y, δ=3, x₀=21) = max(zero(x), (x - x₀ + δ) / δ)
17 Fu = Relaxation(target=U; mask, rate=1)
18 Fv = Relaxation(target=0; mask, rate=1)
19
20 model = NonhydrostaticModel(; grid, closure, boundary_conditions, forcing=(u=Fu, v=Fv))

```

Listing 4: Direct numerical simulation of flow past a cylinder at various Reynolds numbers Re . The domain is periodic in x and a sponge layer on the right side of relaxes the solution to $\mathbf{u} = u_\infty \hat{\mathbf{x}}$ with $u_\infty = 1$. The experiment can be converted to a large eddy simulation (thereby sending $Re \rightarrow \infty$) by replacing the no-slip boundary conditions with an appropriate drag model and either (i) using a turbulence closure like AnisotropicMinimumDissipation or SmagorinskyLilly or (ii) using the WENO(order=9) advection scheme with no turbulence closure. Visualizations of the DNS and LES cases are shown in figure 2.

350 FluxBoundaryCondition (direct imposition of fluxes), and OpenBoundaryCondition (for
351 non-trivial boundary-normal velocity fields).

352 Results obtained with listing 4 for $Re = 100$, $Re = 1000$, and a modified version of
353 listing 4 for large eddy simulation ($Re \rightarrow \infty$) are visualized in figure 2. To adapt listing 4
354 for LES, the closure is eliminated in favor of a 9th-order WENO advection scheme, and
355 the no-slip boundary condition is replaced with a quadratic drag boundary condition with
356 a drag coefficient estimated from similarity theory using a constant estimated roughness
357 length.

358 3.1.2 The NonhydrostaticModel tracer conservation equation

359 The buoyancy term in (6) requires tracers, and can be formulated to use buoyancy itself
360 as a tracer, or to depend on temperature T and salinity S . For seawater, a 54-term polynomial
361 approximation TEOS10EquationOfState (McDougall & Barker, 2011; Roquet, Madec,
362 McDougall, & Barker, 2015) is implemented in the auxiliary package SeawaterPolynomials,
363 along with quadratic approximations to TEOS-10 (Roquet, Madec, Brodeau, & Nycander,
364 2015) and a LinearEquationOfState. All tracers — either “active” tracers required to
365 compute the buoyancy term, as well as additional user-defined passive tracers — obey the tracer
366 conservation equation

$$\partial_t c = -\underbrace{(\mathbf{u} \cdot \nabla) c - (\mathbf{u}_g \cdot \nabla) c}_{\text{advection}} - \underbrace{\nabla \cdot \mathbf{J}_c}_{\text{closure}} + \underbrace{S_c}_{\text{biogeochemistry}} + \underbrace{F_c}_{\text{forcing}}, \quad (9)$$

367 where c represents any tracer, c_g represents a prescribed background tracer concentration
368 for c , \mathbf{J}_c is a tracer flux associated with molecular diffusion or subgrid turbulence, S_c is a
369 source or sink term associated with biogeochemical transformations, and F_c is a user-defined
370 source or sink.

371 A simulation with a passive tracer having a user-defined source term is illustrated by
372 listing 2 and figure 1. For a second example, we consider freshwater cabbeling. Cabbeling

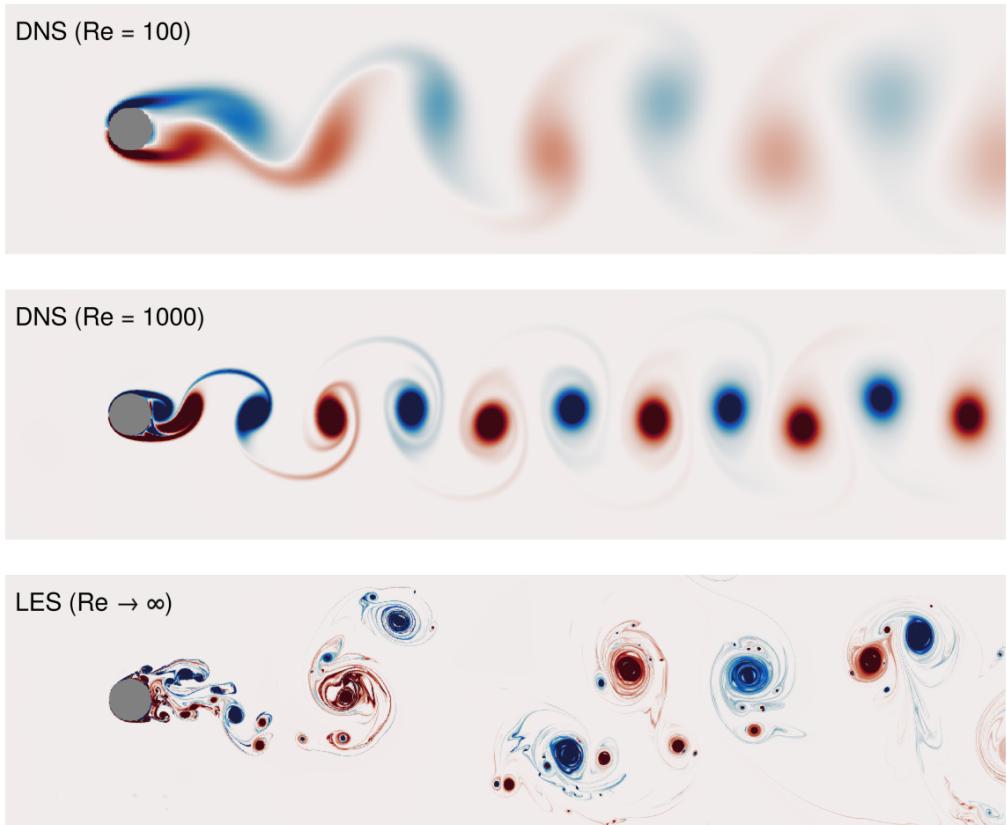


Figure 2: Vorticity snapshots in simulations of flow around a cylinder. The top two panels show vorticity in direct numerical simulations (DNS) that use a molecular ScalarDiffusivity closure and Centered(order=2) advection. The bottom panel shows a large eddy simulation (LES) with no closure and a WENO(order=9) advection scheme.

occurs when two water masses of similar density mix to form a new water mass which,
 due to the nonlinearity of the equation of state, is denser than either of its constituents.
 Freshwater, for example, is densest at 4 degrees Celsius, while 1- and 7.55-degree water are
 lighter with roughly the same density. We implement a direct numerical simulation in which
 7.55-degree water overlies 1-degree water, using the TEOS10EquationOfState provided by
 the auxiliary package SeawaterPolynomials. The script is shown in listing 5. The resulting
 density and temperature fields after 1 minute of simulation are shown in figure 3. Note that
 the TEOS10EquationOfState typically depends on both temperature and salinity tracers,
 but listing 5 specifies a constant salinity $S = 0$ and thus avoids allocating memory for or
 simulating salinity directly.

```

383
384     1 grid = RectilinearGrid(GPU()), topology = (Bounded, Flat, Bounded),
385             size = (4096, 1024), x = (0, 2), z = (-0.5, 0))
386
387     2
388     3 closure = ScalarDiffusivity(v=1.15e-6, κ=1e-7)
389
390     4
391     5 using SeawaterPolynomials: TEOS10EquationOfState
392     6 equation_of_state = TEOS10EquationOfState(reference_density=1000)
393
394     7 buoyancy = SeawaterBuoyancy(gravitational_acceleration = 9.81);
395             constant_salinity = 0, # set S=0 and simulate T only
396             equation_of_state)
397
398     8
399     9 model = NonhydrostaticModel(; grid, buoyancy, closure, tracers=:T)
400
401     10
402     11
403     12
404     13
405     14
406     15 Ti(x, z) = z > -0.25 ? 7.55 : 1
407     16 Ξi(x, z) = 1e-2 * randn()
408     17 set!(model, T=Ti, u=Ξi, v=Ξi, w=Ξi)

```

Listing 5: Direct numerical simulation of convective turbulence driven by cabbeling between 1- and 7.55-degree freshwater. ν denotes viscosity and κ denotes the tracer diffusivity. The diffusivity may also be set independently for each tracer.

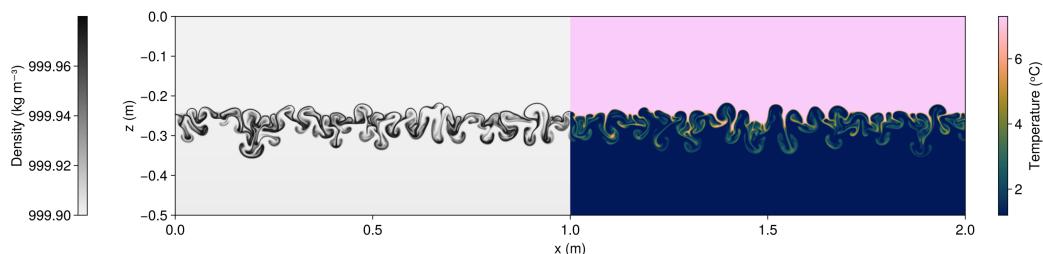


Figure 3: Density and temperature at $t = 1$ minute in a direct numerical simulation of cabelling in freshwater. Note that both fields span from $x = 0$ to $x = 2$ meters; only the left half of the density field and the right half of the temperature field are shown.

We next consider a large eddy simulation of the “Eady problem” (Eady, 1949). In the Eady problem, perturbations evolve around a basic state with constant shear Λ in thermal wind balance with a constant meridional buoyancy gradient $f\Lambda$, such that

$$u = \underbrace{\Lambda z}_{\stackrel{\text{def}}{=} U} + u', \quad \text{and} \quad b = \underbrace{-f\Lambda y}_{\stackrel{\text{def}}{=} B} + b'. \quad (10)$$

We use Oceananigans' `BackgroundFields` to simulate the nonlinear evolution of (u', v, w) and b' expanded around U and B in a doubly-periodic domain. We impose an initially stable density stratification with $b' = N^2 z$ and $N^2 = 10^{-7} \text{ s}^{-2}$ superposed with random noise. The Richardson number of the initial condition is $Ri = N^2 / \partial_z U = N^2 / \Lambda$; we choose

mean shear Λ so that $Ri = 1$, which guarantees the basic is unstable to baroclinic instability but stable to symmetric and Kelvin-Helmholtz instability (Stone, 1971). A portion of the script is shown in listing 6.

```

414
415     1 grid = RectilinearGrid(GPU(); size = (1024, 1024, 64),
416           x = (0, 4096), y = (0, 4096), z = (0, 128),
417           topology=(Periodic, Periodic, Bounded))
418
419     2 f, N2, Ri = 1e-4, 1e-7, 1
420     3 parameters = (f=f, Λ=sqrt(N2/Ri)) # U = Λz, so Ri = N2 / ∂z(U) = N2 / Λ and Λ = N / √Ri.
421
422     4 @inline U(x, y, z, t, p) = + p.Λ * z
423     5 @inline B(x, y, z, t, p) = - p.f * p.Λ * y
424
425     6 background_fields = (u = BackgroundField(U; parameters),
426           b = BackgroundField(B; parameters))
427
428     7 model = NonhydrostaticModel(; grid, background_fields,
429           advection = WENO(order=9), coriolis = FPlane(; f),
430           tracers = :b, buoyancy = BuoyancyTracer())
431
432     8 Δz = minimum_zspacing(grid)
433     9 b1(x, y, z) = N2 * z + 1e-2 * N2 * Δz * (2rand() - 1)
434    10 set!(model, b=b1)
435
436

```

Listing 6: Large eddy simulation of the “Eady problem” (cite something) expanded around the background geostrophic shear with $Ri = 1$.

Our Eady simulation uses fully-turbulence-resolving resolution with 4 meter horizontal spacing and 2 meter vertical spacing in a $4\text{ km} \times 4\text{ km} \times 128\text{ m}$ domain and runs for 30 days on a single Nvidia H100 GPU. Four snapshots of vertical vorticity normalized by f (the Rossby number) are shown in figure 4, illustrating the growth of kilometer-scale vortex motions amid bursts of meter-scale three-dimensional turbulence that develop along thin filaments of vertical vorticity and vertical shear. This simple configuration captures a competition between baroclinic instability, which acts to “restratify” or strengthen boundary layer stratification, and three-dimensional turbulent mixing driven either by a forward cascade from kilometer-scale motions (Molemaker et al., 2010; Dong et al., 2024) or atmospheric storms (Boccaletti et al., 2007; Callies & Ferrari, 2018).

Finally, we illustrate Oceananigans’ capabilities for realistic, three-dimensional large eddy simulations in complex geometries by simulating temperature- and salinity-stratified tidal flow past a headland, reminiscent of an extensively observed and modeled flow past Three Tree Point in Puget Sound in the Pacific Northwest of the United States (Pawlak et al., 2003; Warner & MacCready, 2014). The bathymetry involves a sloping wedge that juts from a square-sided channel, such that

$$z_b(x, y) = -H \left(1 + \frac{y + |x|}{\delta} \right), \quad (11)$$

where L is the half-channel width (the total width is $2L$), $\delta = L/2$ represents the scale of the bathymetry, $H = 128\text{ m}$ is the depth of the channel, and $z = z_b(x, y)$ is the height of the bottom. The flow is driven by a tidally-oscillating boundary velocity

$$U(t) = U_2 \sin \left(\frac{2\pi t}{T_2} \right) \quad (12)$$

imposed at the east and west boundaries. Here, $T_2 = 12.421$ hours is the period of the semi-diurnal lunar tide, and $U_2 = 0.15\text{ m s}^{-1}$ is the characteristic tidal velocity around Three Tree Point. The initial temperature and salinity are

$$T|_{t=0} = 12 + 4 \frac{z}{H} ^\circ\text{C}, \quad \text{and} \quad S|_{t=0} = 32 \text{ g kg}^{-1}. \quad (13)$$

A portion of the script that implements this simulation is shown in listing 7.

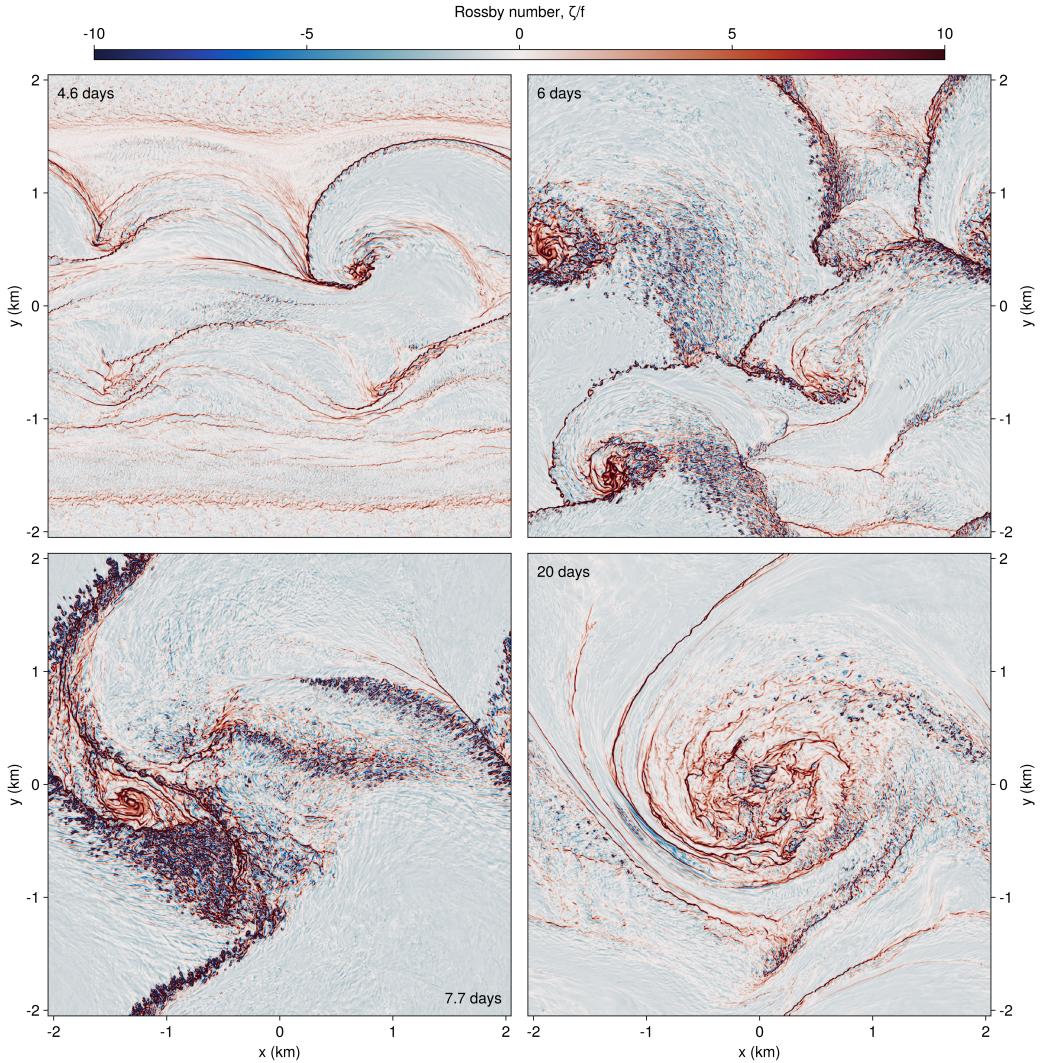


Figure 4: Surface vertical vorticity in a large eddy simulation of the Eady problem with $Ri = 1$ initially, after $t = 4.6, 6, 7.7$, and 20 days. The grid spacing is $4 \times 4 \times 2$ meters in x, y, z . Part of the script that produces this simulation is shown in listing 6.

```

461
462     1 H, L = 256meters, 1024meters
463     2 δ = L / 2
464     3 x, y, z = (-3L, 3L), (-L, L), (-H, 0)
465     4 Nz = 64
466
467     5
468     6 grid = RectilinearGrid(GPU()); size=(6Nz, 2Nz, Nz), halo=(6, 6, 6),
469           x, y, z, topology=(Bounded, Bounded, Bounded))
470     7
471     8
472     9 wedge(x, y) = -H * (1 + (y + abs(x)) / 8)
473    10 grid = ImmersedBoundaryGrid(grid, GridFittedBottom(wedge))
474    11
475    12 T₂ = 12.421hours
476    13 U₂ = 0.1 # m/s
477    14
478    15 @inline Fu(x, y, z, t, p) = 2π * p.U₂ / p.T₂ * cos(2π * t / p.T₂)
479    16 @inline U(x, y, z, t, p) = p.U₂ * sin(2π * t / p.T₂)
480    17 @inline U(y, z, t, p) = U(zero(y), y, z, t, p)
481
482    18 open_bc = PerturbationAdvectionOpenBoundaryCondition(U; inflow_timescale = 2minutes,
483                               outflow_timescale = 2minutes,
484                               parameters=(; U₂, T₂))
485
486    19 u_bcs = FieldBoundaryConditions(east = open_bc, west = open_bc)
487
488    20
489    21
490    22
491    23 u_bcs = FieldBoundaryConditions(east = open_bc, west = open_bc)
492
493    24
494    25 @inline ambient_temperature(x, z, t, H) = 12 + 4z/H
495    26 @inline ambient_temperature(x, y, z, t, H) = ambient_temperature(x, z, t, H)
496    27 ambient_temperature_bc = ValueBoundaryCondition(ambient_temperature; parameters = H)
497    28 T_bcs = FieldBoundaryConditions(east = ambient_temperature_bc,
498                                         west = ambient_temperature_bc)
499
500    30
501    31 ambient_salinity_bc = ValueBoundaryCondition(32)
502    32 S_bcs = FieldBoundaryConditions(east = ambient_salinity_bc, west = ambient_salinity_bc)
503
504    33
505    34 model = NonhydrostaticModel(; grid, tracers = (:T, :S),
506                                     buoyancy = SeawaterBuoyancy(equation_of_state=
507                                         TEOS10EquationOfState())),
508                                     advection = WENO(order=9),
509                                     coriolis = FPlane(latitude=47.5),
510                                     boundary_conditions = (; T=T_bcs, u = u_bcs, S = S_bcs))
511
512    39
513    40 Tᵢ(x, y, z) = ambient_temperature(x, y, z, 0, H)
514
515    41
516    42 set!(model, T=Tᵢ, S=32, u=U(0, 0, 0, 0, (; U₂, T₂)))

```

Listing 7: Large eddy simulation of flow past a headland reminiscent of Three Tree Point in the Pacific Northwest (see Pawlak et al., 2003; Warner & MacCready, 2014).

The oscillatory, turbulent flow is visualized in figure 5, and the calculation of the Ertel Potential Vorticity seen in figure 5c is done with the companion package Oceanostics.

3.2 Hydrostatic model with a free surface

The HydrostaticFreeSurfaceModel solves the *hydrostatic*, rotating Boussinesq equations with a free surface. The hydrostatic approximation, inherent to the HydrostaticFreeSurfaceModel, means that the vertical momentum equation is replaced by an integral for the hydrostatic pressure and that the vertical velocity is diagnosed from the continuity equation. The HydrostaticFreeSurfaceModel therefore does not need a three-dimensional Poisson equation for pressure. The numerical algorithms and computational performance of the HydrostaticFreeSurfaceModel are described in more detail by ([Silvestri, Wagner, Constantinou, et al., 2024](#)).

In the HydrostaticFreeSurfaceModel, the horizontal momentum $\mathbf{u}_h = u\hat{\mathbf{x}} + v\hat{\mathbf{y}}$ evolves according to

$$\partial_t \mathbf{u}_h = -\nabla_h p - \underbrace{g \nabla_h \eta}_{\text{free surface}} - \underbrace{(\mathbf{u} \cdot \nabla) \mathbf{u}_h}_{\text{momentum advection}} - \underbrace{\mathbf{f} \times \mathbf{u}}_{\text{Coriolis}} - \underbrace{\nabla \cdot \boldsymbol{\tau}}_{\text{closure}} + \underbrace{\mathbf{F}_{uh}}_{\text{forcing}}, \quad (14)$$

where p is the hydrostatic pressure anomaly, η is the free surface displacement, $\mathbf{u} = u\hat{\mathbf{x}} + v\hat{\mathbf{y}} + w\hat{\mathbf{z}}$ is the three-dimensional velocity, \mathbf{f} is the background vorticity associated with

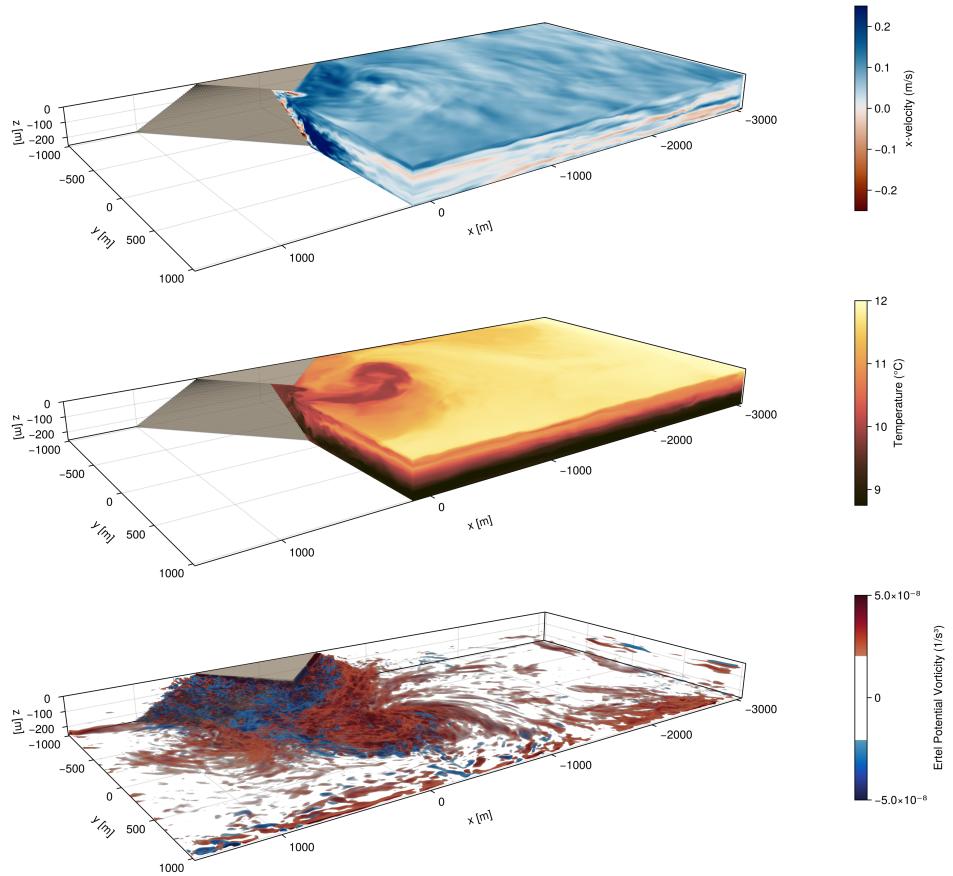


Figure 5: Along-channel velocity, temperature, and Ertel potential vorticity in a tidally-forced flow past an idealized headland with open boundaries. The tidal flow occurs in the x -directions and the snapshot depicts the flow just after the tide has turned to the negative- x direction.

523 a rotating frame of reference, $\boldsymbol{\tau}$ is the stress associated with subgrid turbulent horizontal
 524 momentum transport, and \mathbf{F}_{uh} is a body force. Horizontal momentum advection can be
 525 formulated in three ways,

$$(u \cdot \nabla) \cdot \mathbf{u}_h = \begin{cases} \nabla \cdot (\mathbf{u} \mathbf{u}_h) & \text{"flux form"}, \\ \zeta \hat{\mathbf{z}} \times \mathbf{u}_h + w \partial_z \mathbf{u}_h + \nabla_h \frac{1}{2} |\mathbf{u}_h|^2 & \text{VectorInvariant}, \\ \zeta \hat{\mathbf{z}} \times \mathbf{u}_h - \mathbf{u}_h \partial_z w + \partial_z (w \mathbf{u}_h) + \nabla_h \frac{1}{2} |\mathbf{u}_h|^2 & \text{WENOVectorInvariant}, \end{cases} \quad (15)$$

526 where the "flux form" treats momentum advection in the same way as for the NonhydrostaticModel.
 527 The numerical implementation of the WENOVectorInvariant formulation, which
 528 leverages Weighted Essentially Non-Oscillatory (WENO) reconstructions to selectively and
 529 minimally dissipate enstrophy and the variance of divergence (see section 4), is described
 530 by Silvestri, Wagner, Campin, et al. (2024).

531 The hydrostatic pressure anomaly is determined diagnostically from hydrostatic balance,
 532

$$\partial_z p = b, \quad (16)$$

533 which replaces the prognostic vertical momentum equation $\hat{\mathbf{z}} \cdot (6)$ used by the Nonhydro-
 534 staticModel. The vertical velocity is obtained diagnostically from the continuity equation,

$$\partial_z w = -\nabla_h \cdot \mathbf{u}_h, \quad (17)$$

535 and the free surface displacement η obeys the linearized equation

$$\partial_t \eta = w|_{z=0}. \quad (18)$$

536 Tracer evolution is governed by the conservation law

$$\partial_t c = -\underbrace{(\mathbf{u} \cdot \nabla) c}_{\text{tracer advection}} - \underbrace{\nabla \cdot \mathbf{J}_c}_{\text{closure}} + \underbrace{S_c}_{\text{biogeochemistry}} + \underbrace{F_c}_{\text{forcing}}, \quad (19)$$

537 which is identical to NonhydrostaticModel except that background fields are not supported.
 538 Additionally, the velocity field \mathbf{u} can be prescribed rather than evolved.

539 Listing 8 implements a simulation of tidally-forced stratified flow over a series of
 540 randomly-positioned Gaussian seamounts.

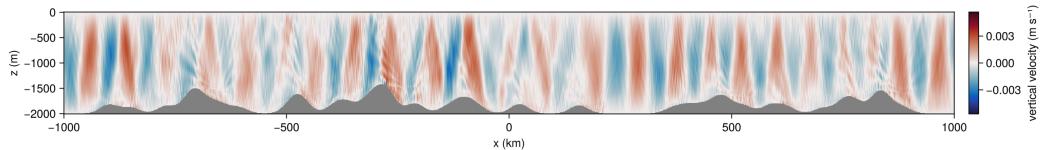


Figure 6: Vertical velocity of an internal wave field excited by tidally-forced stratified flow over superposition of randomly-positioned Gaussian seamounts, after 16 tidal periods.

541 3.2.1 Vertical mixing parameterizations

542 Oceananigans' vertical mixing parameterizations are closures that predict the vertical
 543 fluxes of tracers and momentum. Depending on the parameterization, the evolution of
 544 auxiliary tracers like turbulent kinetic energy and the turbulent kinetic energy dissipation
 545 rate may also be simulated. Vertical mixing parameterizations are useful for hydrostatic
 546 simulations where vertical mixing is otherwise unresolved due to a coarse horizontal grid

```

1 using Oceananigans, Oceananigans.Units
2
3 grid = RectilinearGrid(size = (2000, 200),
4                         x = (-1000kilometers, 1000kilometers),
5                         z = (-2kilometers, 0),
6                         halo = (4, 4),
7                         topology = (Periodic, Flat, Bounded))
8
9 h₀ = 100           # typical mountain height (m)
10 δ = 20kilometers # mountain width (m)
11 seamounts = 42
12 W = grid.Lx - 4δ
13 x₀ = W .* (rand(seamounts) .- 1/2) # mountains' positions ∈ [-Lx/2+2δ, Lx/2-2δ]
14 h = h₀ .* (1 .+ rand(seamounts)) # mountains' heights ∈ [h₀, 2h₀]
15
16 bottom(x) = -grid.Lz + sum(h[s] * exp(-(x - x₀[s])^2 / 2δ^2) for s = 1:seamounts
17 grid = ImmersedBoundaryGrid(grid, GridFittedBottom(bottom))
18
19 T₂ = 12.421hours # period of M₂ tide constituent
20 @inline tidal_forcing(x, z, t, p) = p.U₂ * 2π / p.T₂ * sin(2π / p.T₂ * t)
21 u_forcing = Forcing(tidal_forcing, parameters=(; U₂=0.1, T₂=T₂))
22
23 model = HydrostaticFreeSurfaceModel(; grid, tracers=:b, buoyancy=BuoyancyTracer(),
24                                     momentum_advection = WENO(),
25                                     tracer_advection = WENO(),
26                                     forcing = (; u = u_forcing))
27
28 b₁(x, z) = 1e-5 * z
29 set!(model, b=b₁)

```

Listing 8: Two-dimensional simulation of tidally-forced stratified flow over a superposition of randomly-positioned Gaussian seamounts.

spacing. For example, such regional and global configurations, horizontal grid spacing typically varies from $O(100\text{ m})$ to $O(100\text{ km})$. Listing 9 implements a simulation of wind-driven vertical mixing in a single column model using two parameterizations: CATKE (Wagner et al., 2024), which has one additional equation for the evolution of turbulent kinetic energy (TKE), and k - ϵ (Umlauf & Burchard, 2005), which has two additional equations for TKE and the TKE dissipation rate. Figure 7 plots the result.

3.3 Global ocean simulations with ClimaOcean

The HydrostaticFreeSurfaceModel can be used to simulate regional or global ocean circulation. To illustrate global simulation with the HydrostaticFreeSurfaceModel, we implement a global simulation using ClimaOcean, which computes surface fluxes between a prescribed atmosphere and a hydrostatic ocean simulation implemented using Oceananigans. ClimaOcean additionally provides utilities for downloading and interfacing with JRA55 reanalysis data (Tsujino et al., 2018), jbuilding grids based on Earth bathymetry and initializing simulations from the ECCO state estimate (Forget et al., 2015). Code for a 1/12th degree simulation distributed over 8 GPUs is given in listing 10. The surface speed after 180 days of simulation time is shown in figure 8. For more information about Oceananigans GPU performance in global configurations see (Silvestri, Wagner, Constantinou, et al., 2024).

```

565
566
567 using ClimaOcean, Oceananigans, Oceananigans.Units, Dates, CFTime
568
569 # 1/12th degree
570 Nx = 4320
571 Ny = 1800
572 Nz = 40
573 z_faces = exponential_z_faces(; Nz, depth=6000)
574 partition = Partition(8) # Distribute simulation across 8 GPUs

```

```

1 using Oceananigans
2 using Oceananigans.Units
3
4 function vertical_mixing_simulation(closure; N2=1e-5, Jb=1e-7, tx=-5e-4)
5   grid = RectilinearGrid(size=50, z=(-200, 0), topology=(Flat, Flat, Bounded))
6   buoyancy = BuoyancyTracer()
7
8   b_bcs = FieldBoundaryConditions(top=FluxBoundaryCondition(Jb))
9   u_bcs = FieldBoundaryConditions(top=FluxBoundaryCondition(tx))
10
11  if closure isa CATKEVerticalDiffusivity
12    tracers = (:b, :e)
13  elseif closure isa TKEDissipationVerticalDiffusivity
14    tracers = (:b, :e, :ε)
15  end
16
17  model = HydrostaticFreeSurfaceModel(; grid, closure, tracers, buoyancy,
18                                     boundary_conditions=(u=u_bcs, b=b_bcs))
19
20  bi(z) = N2 * z
21  set!(model, b=bi)
22
23  simulation = Simulation(model, Δt=1minute, stop_time=24hours)
24  return run!(simulation)
25 end

```

Listing 9: Comparison of two vertical mixing parameterizations in the evolution of an initially linearly stratified boundary layer subjected to stationary surface fluxes of buoyancy and momentum.

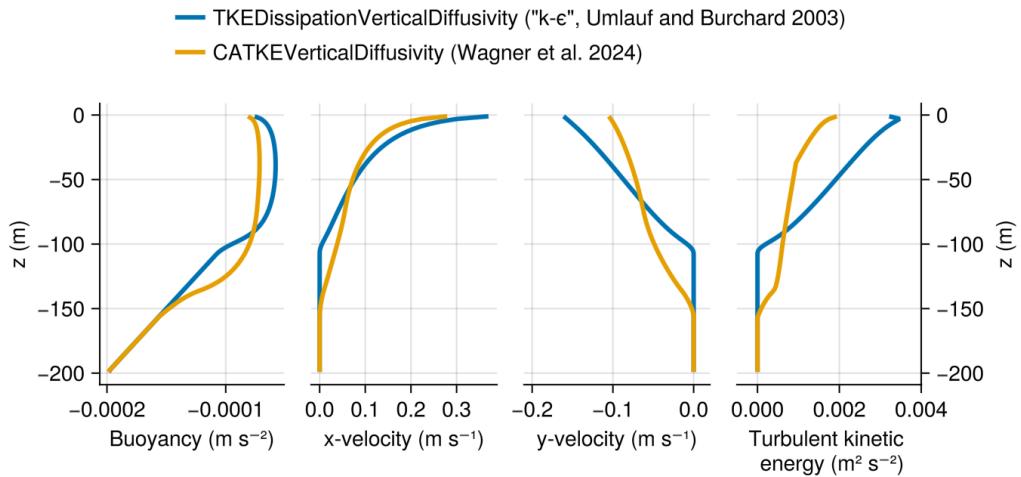


Figure 7: Vertical mixing parameterizations.

```

575     9 arch = Distributed(GPU(); partition)
576    10 grid = LatitudeLongitudeGrid(arch; size=(Nx, Ny, Nz), halo=(7, 7, 7),
577                                     longitude=(0, 360), latitude=(-75, 75), z=z_faces)
578    11
579    12 bathymetry = ClimaOcean.regrid_bathymetry(grid) # based on ETOPO1
580    13 grid = ImmersedBoundaryGrid(grid, GridFittedBottom(bathymetry))
581    14
582    15 # Build an ocean simulation initialized to the ECCO state estimate on Jan 1, 1993
583    16 ocean = ClimaOcean.ocean_simulation(grid)
584    17 date = DateTimeProlepticGregorian(1993, 1, 1)
585    18 set!(ocean.model, T = ClimaOcean.ECCOMetadata(:temperature; date),
586          S = ClimaOcean.ECCOMetadata(:salinity; date))
587    19
588    20 # Global ocean simulation (no sea ice) forced by JRA55 reanalysis
589    21 backend = JRA55NetCDFBackend(41))
590    22 atmosphere = ClimaOcean.JRA55_prescribed_atmosphere(arch; backend)
591    23 coupled_model = ClimaOcean.OceanSeaIceModel(ocean; atmosphere)
592    24 simulation = Simulation(coupled_model, Δt=5minutes, stop_time=180days)
593    25 run!(simulation)
594
595

```

Listing 10: A near-global simulation on a LatitudeLongitudeGrid distributed across 4 GPUs, leveraging ClimaOcean.

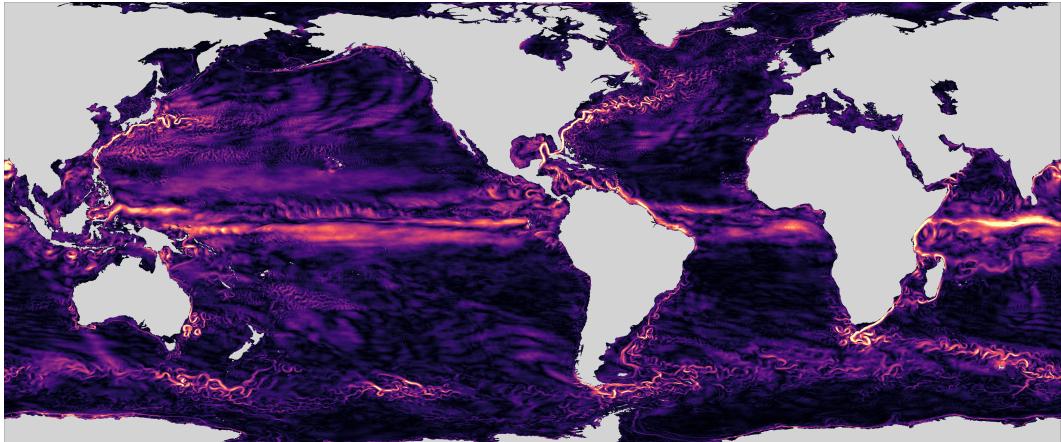


Figure 8: Surface speed in a near-global ocean simulation at 1/12th degree forced by JRA55 atmospheric reanalysis and initialized from the ECCO state estimate.

4 Finite volume spatial discretization

Oceananigans uses a finite volume method in which fields are represented discretely by their average value over small local regions or “finite volumes” of the domain. Listing 11 discretizes $c = e^x y$ on three different grids that cover the unit square.

```

600
601     1 topology = (Bounded, Bounded, Flat)
602     2 x = y = (0, 1)
603     3 c(x, y) = exp(x) * y
604     4
605     5 fine_grid = RectilinearGrid(size=(1024, 1024); x, y, topology)
606     6 c_fine = CenterField(fine_grid)
607     7 set!(c_fine, c)
608     8
609     9 medium_grid = RectilinearGrid(size=(16, 16); x, y, topology)
610    10 c_medium = CenterField(medium_grid)
611    11 regrid!(c_medium, c_fine)
612    12
613    13 coarse_grid = RectilinearGrid(size=(4, 4); x, y, topology)
614    14 c_coarse = CenterField(coarse_grid)
615    15 regrid!(c_coarse, c_medium)
616
617

```

618

Listing 11: Finite volume discretization of $e^x y$ on three grids over the unit square. The fields are visualized in figure 9. The meaning of the “Center” in “CenterField” is discussed below.

At the finest resolution, each cell-averaged value c_{ij}^{fine} is computed approximately using set! to evaluate $e^x y$ at the center of each finite volume, where i, j denote the x and y indices of the finite volumes. At medium and coarse resolution, the c_{ij}^{medium} and c_{ij}^{coarse} are computed by averaging or “regridding” fields discretized at a higher resolution. This computation produces three fields with identical integrals over the unit square. For example, integrals are computed exactly by summing discrete fields over all cells,

$$\int c \, dx \, dy = \sum_{i,j}^{1024,1024} \mathcal{V}_{ij}^{\text{fine}} c_{ij}^{\text{fine}} = \sum_{i,j}^{16,16} \mathcal{V}_{ij}^{\text{medium}} c_{ij}^{\text{medium}} = \sum_{i,j}^{4,4} \mathcal{V}_{ij}^{\text{coarse}} c_{ij}^{\text{coarse}}, \quad (20)$$

where \mathcal{V}_{ij} is the “volume” of the cell with indices i, j (more accurately an “area” in this two-dimensional situation). Figure 9 visualizes the three fields.

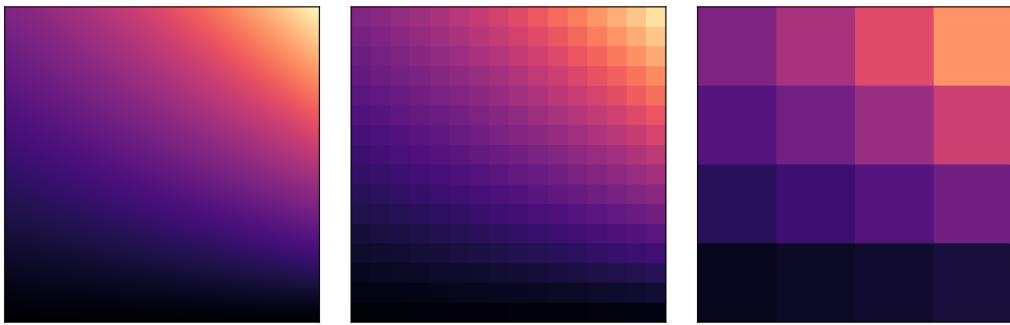


Figure 9: Finite volume discretization of $e^x y$ on the unit square at three different resolutions.

The discrete calculus and arithmetic operations required to solve the governing equations of the NonhydrostaticModel and HydrostaticFreeSurfaceModel use the system of “staggered grids” described by (Arakawa, 1977). Both models use “C-grid” staggering, where cells for tracers, pressure, and the divergence of the velocity field $\nabla \cdot \mathbf{u}$ are co-located, and cells for velocity components $\mathbf{u} = (u, v, w)$ are staggered by half a cell width in the x -, y -, and z -direction, respectively. Listing 12 illustrates grid construction and notation for a one-dimensional staggered grid with unevenly-spaced cells. Figure 10 visualizes 2- and 3-dimensional staggered grids, indicating the location of certain variables.

635
636
637
638
639
640
641
642
643
644
645
646
647

```
1 using Oceananigans
2
3 grid = RectilinearGrid(topology=(Bounded, Flat, Flat), size=4, x=[0, 0.2, 0.3, 0.7, 1])
4
5 u = Field{Face, Center, Center}(grid)
6 c = Field{Center, Center, Center}(grid)
7
8 xnodes(u)      # [0.0, 0.2, 0.3, 0.7, 1.0]
9 xnodes(c)      # [0.1, 0.25, 0.5, 0.85]
10 location(∂x(c)) # (Face, Center, Center)
```

Listing 12: A one-dimensional staggered grid.

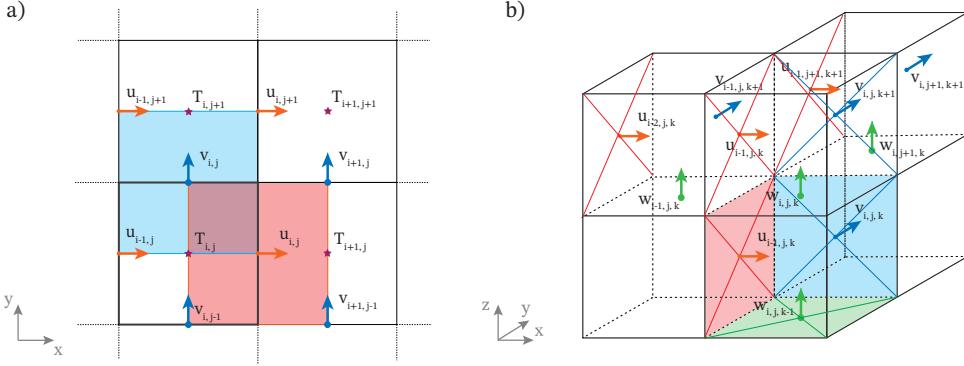


Figure 10: Locations of cell centers and interfaces on a two-dimensional (a) and three-dimensional (b) staggered grid. In (a), the red and blue shaded regions highlight the volumes in the dual u -grid and v -grid, located at (Face, Center, Center) and (Center, Face, Center), respectively. In (b), the shaded regions highlight the facial areas used in the fluxes computations, denoted with \mathcal{A}_x , \mathcal{A}_y , and \mathcal{A}_z .

649 4.1 A system of composable operators

650 A convention for indexing is associated with staggered locations. Face indices are “left”
 651 of cell indices. This means that difference operators acting on fields at cells differ from
 652 those that act on face fields. To illustrate this we introduce Oceananigans-like difference
 653 operators,

```
654 1 δxfcc(i, j, k, grid, c) = c[i, j, k] - c[i-1, j, k]
655 2 δxccc(i, j, k, grid, u) = u[i+1, j, k] - u[i, j, k]
```

660 where superscripts denote the location of the *result* of the operation. For example, the dif-
 661 ference δ_x^{fcc} acts on fields located at ccc (meaning cell Center in the x , y and z directions
 662 respectively). Complementary to the difference operators are reconstruction of “interpo-
 663 lation” operators,

```
664 1 ℑxfcc(i, j, k, grid, c) = (c[i, j, k] + c[i-1, j, k]) / 2
665 2 ℑxccc(i, j, k, grid, u) = (u[i+1, j, k] + u[i, j, k]) / 2
```

670 The prefix arguments i , j , k , $grid$ are more than convention: the prefix enables
 671 system for *composing* operators. For example, defining

```
672 1 δxfcc(i, j, k, grid, f::Function, args...) =
673 2     f(i, j, k, grid, args...) - f(i-1, j, k, grid, args...)
674 3
675 4 δxccc(i, j, k, grid, f::Function, args...) =
676 5     f(i+1, j, k, grid, args...) - f(i, j, k, grid, args...)
```

681 leads to a concise definition of the second-difference operator:

```
682 1 δ2xccc(i, j, k, grid, f::Function, a...) = δxccc(i, j, k, grid, δxfcc, f, a...)
```

687 Operator composition is used throughout Oceanangians source code to implement stencil
 688 operations.

689 4.2 Tracer flux divergences, advection schemes, and reconstruction

690 The divergence of a tracer flux $\mathbf{J} = J_x \hat{\mathbf{x}} + J_y \hat{\mathbf{y}} + J_z \hat{\mathbf{z}}$ is discretized conservatively by
 691 the finite volume method via

$$\nabla \cdot \mathbf{J} \approx \frac{1}{V_c} \left[\underbrace{\delta_x(\mathcal{A}_x J_x)}_{\text{fcc}} + \underbrace{\delta_y(\mathcal{A}_y J_y)}_{\text{cfc}} + \underbrace{\delta_z(\mathcal{A}_z J_z)}_{\text{ccf}} \right], \quad (21)$$

692 where $\delta_x, \delta_y, \delta_z$ are difference operators in x, y, z , V_c denotes the volume of the tracer cells,
 693 $\mathcal{A}_x, \mathcal{A}_y$, and \mathcal{A}_z denote the areas of the tracer cell faces with surface normals $\hat{\mathbf{x}}, \hat{\mathbf{y}}$, and
 694 $\hat{\mathbf{z}}$, respectively. Equation (21) indicates the location of each flux component: fluxes into
 695 tracers cell at ccc are computed at the cell faces located at fcc, cfc, and ccf.

696 The advective tracer flux in (9) is written in “conservative form” using incompressibility
 697 (2), and then discretized similarly to (21) to form

$$\mathbf{u} \cdot \nabla c = \nabla \cdot (\mathbf{u}c) \approx \frac{1}{V_c} \left[\delta_x(\mathcal{A}_x u|c|_x) + \delta_y(\mathcal{A}_y v|c|_y) + \delta_z(\mathcal{A}_z w|c|_z) \right], \quad (22)$$

698 where $|c|_x$ denotes a *reconstruction* of c in the x -direction from its native location ccc to the
 699 tracer cell interface at fcc; $|c|_y$ and $|c|_z$ in (22) are defined similarly.

700 The advective fluxes $\mathbf{u}c$ must be computed on interfaces between tracer cells, where
 701 the approximate value of c must be reconstructed. (Velocity components like u must also
 702 be reconstructed on interfaces. Within the C-grid framework, we approximate u on tracer
 703 cell interfaces directly using the values u_{ijk} , which represent u averaged over a region en-
 704 compassing the interface.) The simplest kind of reconstruction is Centered(order=2), which
 705 is equivalent to taking the average between adjacent cells,

$$\langle c \rangle_i = \frac{1}{2} (c_i + c_{i-1}), \quad (23)$$

706 where $\langle c \rangle_i$ denotes the reconstruction of c on the interface at $x = x_{i-1/2}$. Also in (23)
 707 the j, k indices are implied and we have suppressed the direction x to lighten the notation.
 708 Reconstructions stencils for Center(order=N) are automatically generated for even N up
 709 to $N_{\max} = 12$, where N_{\max} is an adjustable parameter in the source code. All subsequent
 710 reconstructions are described in the x -direction only.

711 Centered schemes should be used when explicit dissipation justified by a *physical* ration-
 712 ale dominates at the grid scale. In scenarios where dissipation is needed solely for artificial
 713 reasons, we find applications for UpwindBiased schemes, which use an odd-order stencil
 714 biased against the direction of flow. For example, UpwindBiased(order=1) and UpwindBi-
 715 ased(order=3) schemes are written

$$u[c]_x^1 = \begin{cases} u c_{i-1} & \text{if } u > 0, \\ u c_i & \text{if } u < 0, \end{cases} \quad \text{and} \quad u[c]_x^3 = \begin{cases} u \frac{1}{6} (-c_{i-2} + 5c_{i-1} + 2c_i) & \text{if } u > 0, \\ u \frac{1}{6} (2c_{i-1} + 5c_i - c_{i+1}) & \text{if } u < 0, \end{cases} \quad (24)$$

716 where $[c]_x^N$ denotes N^{th} -order upwind reconstruction in the x -direction. (Note that $u[c]_x^N = 0$
 717 if $u = 0$.)

718 The compact form of equations (24) demonstrates how upwind schemes introduce vari-
 719 ance dissipation through numerical discretization. In particular, an UpwindBiased(order=1)
 720 reconstruction can be rewritten as a sum of a Centered(order=2) discrete advective flux and
 721 a discrete diffusive flux

$$u[c]_x^1 = u \frac{c_i + c_{i-1}}{2} - \kappa_1 \frac{c_i - c_{i-1}}{\Delta x}, \quad \text{where} \quad \kappa_1 = \frac{|u| \Delta x}{2}. \quad (25)$$

722 Reordering the UpwindBiased(order=3) advective flux in the same manner recovers a sum
 723 of a Centered(order=4) advective flux and a 4th-order hyperdiffusive flux, equivalent to a
 724 finite volume approximation of

$$725 \quad uc + \kappa_3 \frac{\partial^3 c}{\partial x^3}, \quad \text{where } \kappa_3 = \frac{|u|\Delta x^3}{12}. \quad (26)$$

725 UpwindBiased reconstruction can be always reordered to expose a sum of Centered recon-
 726 struction and a high-order diffusive flux with a velocity-dependent diffusivity. The diffusive
 727 operator associated with UpwindBiased(order=1) and UpwindBiased(order=3) is enough to
 728 offset the dispersive errors of the Centered component and, therefore, eliminate the artificial
 729 explicit diffusion needed for stability purposes. However, this approach does not scale to
 730 high order since the diffusive operator associated with a high order UpwindBiased scheme
 731 (5th, 7th, and so on), becomes quickly insufficient to eliminate spurious errors associated
 732 with the Centered component (Godunov, 1959).

733 The inability to achieve high order and, therefore, low dissipation motivated the imple-
 734 mentation of Weighted, Essentially Non-Oscillatory (WENO) reconstruction (C. Shu, 1997;
 735 C.-W. Shu, 2009). WENO is a non-linear reconstruction scheme that combines a set of
 736 odd-order linear reconstructions obtained by stencils that are shifted by a value s relative to
 737 the canonical UpwindBiased stencil, using a weighting scheme for each stencil that depends
 738 on the smoothness of the reconstructed field c . Since the constituent stencils are lower-
 739 order than the WENO order, this strategy yields a scheme whose order of accuracy adapts
 740 depending on the smoothness of the reconstructed field. In smooth regions high-order is
 741 retained, while the order quickly decreases in the presence of noisy regions, decreasing the or-
 742 der of the associated diffusive operator. WENO provies especially useful for high-resolution,
 743 turbulence-resolving simulations (either at meter or planetary scales) without requiring any
 744 additional explicit artificial dissipation (Pressel et al., 2017; Silvestri, Wagner, Campin, et
 745 al., 2024).

746 To illustrate how WENO works we consider a fifth-order WENO scheme for $u > 0$,

$$\{c\}^5 = \gamma_0[c]^{3,0} + \gamma_1[c]^{3,1} + \gamma_2[c]^{3,2}, \quad (27)$$

747 where the notation $[c]^{3,s}$ denotes an UpwindBiased stencil *shifted* by s indices, such that
 748 $[c]^3 \stackrel{\text{def}}{=} [c]^{3,0}$. The shifted upwind stencils $[c]_i^{N,s}$ evaluated at index i are defined

$$749 \quad [c]_i^{3,s} = \frac{1}{6} \begin{cases} -c_{i-1} + 5c_i + 2c_{i+1} & \text{for } s = -1, \\ 2c_{i-2} + 5c_{i-1} - c_i & \text{for } s = 0, \\ 2c_{i-3} - 7c_{i-2} + 11c_{i-1} & \text{for } s = 2. \end{cases} \quad (28)$$

750 The weights $\gamma_s(c)$ are determined by a smoothness metric that produces $\{c\}^5 \approx [c]^5$ when c is
 751 smooth, but limits to the more diffusive $\{c\}^5 \approx [c]^3$ when c changes abruptly. Thus WENO
 752 adaptively introduces dissipation as needed based on the smoothness of c , yielding stable
 753 simulations with a high effective resolution that require no artificial dissipation. WENO
 754 can alternatively be interpreted as adding an implicit hyperviscosity that adapts from low-
 755 to high-order depending on the local nature of the solution. To compute the weights $\gamma_s(c)$,
 we use the WENO-Z formulation (Balsara & Shu, 2000).

756 The properties of Centered, UpwindBiased, and WENO reconstruction are investigated
 757 by listing 13, which simulates the advection of a top hat tracer distribution. The results are
 758 plotted in figure 11.

```
759
760 1 using Oceananigans
761 2
762 3 grid = RectilinearGrid(size=128; x=(-4, 8), halo=6, topology=(Periodic, Flat, Flat))
763 4 advection = WENO(order=9) # Centered(order=2), UpwindBiased(order=3)
764 5 velocities = PrescribedVelocityFields(u=1)
```

```

766   model = HydrostaticFreeSurfaceModel(; grid, velocities, advection, tracers=:c)
767
768   8 top_hat(x) = abs(x) > 1 ? 0 : 1
769   9 set!(model, c = top_hat)
770
771 10 simulation = Simulation(model, Δt=1/grid.Nx, stop_time=4)
772 11 run!(simulation)
773

```

Listing 13: A script that advects a top hat tracer profile in one-dimension with a constant prescribed velocity. We use halo=6 to accommodate schemes up to WENO(order=11).

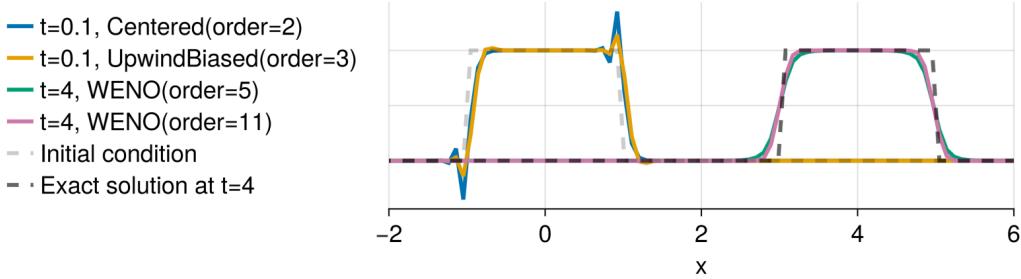


Figure 11: Advection of a top hat tracer distribution in one-dimension using various advection schemes. Centered and Upwind

775 4.2.1 Discretization of momentum advection

The discretization of momentum advection with a flux form similar to (22) is more complex than the tracer case because both the advecting velocity and advected velocity require reconstruction. We use the method described by Ghosh and Baeder (2012) and Pressel et al. (2015), wherein advecting velocities are constructed with a high-order Centered scheme when the advected velocity component is reconstructed with a high-order UpwindBiased or WENO scheme. We have also developed a novel WENO-based method for discretizing momentum advection in the rotational or “vector invariant” form especially appropriate for representing mesoscale and submesoscale turbulent advection on curvilinear grids (Silvestri, Wagner, Campin, et al., 2024).

785 5 Parallelization

Oceananigans supports distributed computations with slab and pencil domain decomposition. The interior domain is extended using “halo” or “ghost” cells that hold the results of interprocessor boundaries. “halo” cells are updated before the computation of tendencies through asynchronous send / receive operations using the message passing interface (MPI) Julia library (Byrne et al., 2021). For a detailed description of the parallelization strategy of the HydrostaticFreeSurfaceModel; see Silvestri, Wagner, Constantinou, et al. (2024). The NonhydrostaticModel implements the same overlap of communication and computation for halo exchange before the calculation of tendencies. For the FFT-based three-dimensional pressure solver, we implement a transpose algorithm that switches between x -local, y -local, and z -local configurations to compute efficiently the discrete transforms. The transpose algorithm for the distributed FFT solver is shown in figure 12.

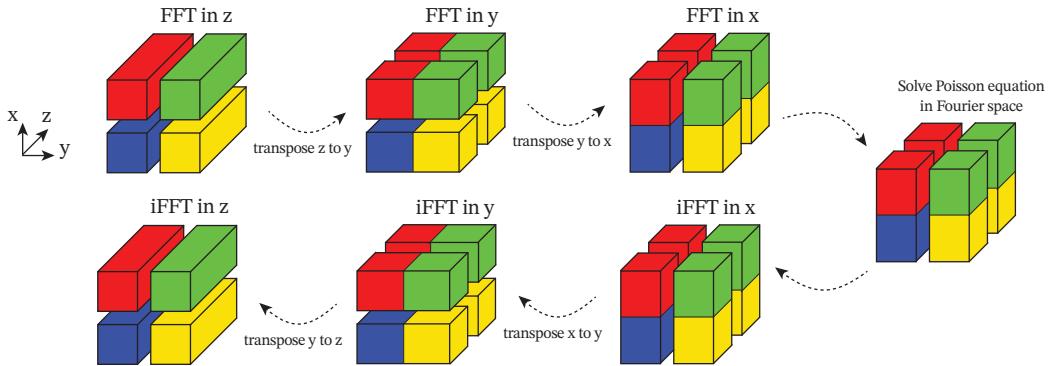


Figure 12: A schematic showing the distributed Poisson solver procedure with a pencil parallelization that divides the domain in two ranks in both x and y . The schematic highlights the data layout in the ranks during each operation.

6 Conclusions

This paper describes GPU-based ocean modeling software called “Oceananigans” written in the high-level Julia programming language. Oceananigans’ exposes ground-breaking performance for simulations of oceanic motion at any scale with an innovative user interface design that makes simple simulations easy and complex, creative simulations possible.

Oceananigans yields a brute force strategy for accuracy: simple, C-grid, WENO numerics for turbulence resolving simulations coupled to the raw power of GPU acceleration. This strategy bucks the trend in dynamical core numerics, especially for oceanography, which in recent years has trended toward unstructured grids, vertically-Lagrangian coordinates. Yet we anticipate that GPUs, driven by machine learning, will only become more powerful. We hope therefore that Oceananigans brute-force strategy will continue to bear fruit. Atmospheric dycores may also benefit from such simplifications, moving away from vertically-Lagrangian coordinates, for example.

Each achievement — ground-breaking performance, physics flexibility, or an innovative design — would, on their own, enable scientific breakthroughs. By assembling these achievements into a single package, however, and eliminating the typical trade-offs between performance, flexibility, and ease-of-use, Oceananigans aims at the even higher goal of accelerating the pace of progress in model development. This matters because ocean modeling software will have to continue to evolve rapidly to keep pace with the advancing state of the field to remain cutting-edge: to continue to use the world’s largest super computers, to continue to present the most productive possible abstractions for both users and developers, to ensure that the pace of parameterization development is as fast as it could be. Right now, ocean model development is arguably too slow.

Appendix A Time stepping and time discretization

In this section we describe time stepping methods and time discretization options for the NonhydrostaticModel and the HydrostaticFreeSurfaceModel.

A1 Time discretization for tracers

Tracers are stepped forward with similar schemes in the NonhydrostaticModel and the HydrostaticFreeSurfaceModel, each of which includes optional implicit treatment of vertical

826 diffusion terms. Equation (9) is abstracted into two components,

$$\partial_t c = G_c + \partial_z (\kappa_z \partial_z c), \quad (\text{A1})$$

827 where, if specified, κ_z is the vertical diffusivity of c to be treated with a VerticallyImplicitTimeDiscretization, and G_c is the remaining component of the tracer tendency from
828 equation 9. (Vertical diffusion treated with an ExplicitTimeDiscretization is also absorbed
829 into G_c .) We apply a semi-implicit time discretization of vertical diffusion to approximate
830 integral of (A1) from t^m to t^{m+1} ,

$$(1 - \Delta t \partial_z \kappa_z^m \partial_z) c^{m+1} = c^m + \int_{t^m}^{t^{m+1}} G_c dt, \quad (\text{A2})$$

832 where $\Delta t \stackrel{\text{def}}{=} t^{m+1} - t^m$. The tendency integral $\int_{t^m}^{t^{m+1}} G_c dt$ is evaluated either using a
833 “quasi”-second order Adams-Bashforth scheme (QAB2, which is actually first-order *lets add a reference*), or a low-storage third-order Runge-Kutta scheme (RK3). For QAB2, the
834 integral in (A2) spans the entire time-step and takes the form
835

$$\frac{1}{\Delta t} \int_{t^m}^{t^{m+1}} G_c dt \approx \left(\frac{3}{2} + \chi\right) G_c^m - \left(\frac{1}{2} + \chi\right) G_c^{m-1}, \quad (\text{A3})$$

836 where χ is a small parameter, chosen by default to be $\chi = 0.1$. QAB2 requires one tendency
837 evaluation per time-step. For RK3, the indices $m = (1, 2, 3)$ correspond to *substages*, and
838 the integral in (A2) takes the form

$$\frac{1}{\Delta t} \int_{t^m}^{t^{m+1}} G_c dt \approx \gamma^m G_c^m - \zeta^m G_c^{m-1}, \quad (\text{A4})$$

839 where $\gamma = (8/15, 5/12, 3/4)$ and $\zeta = (0, 17/60, 5/12)$ for $m = (1, 2, 3)$ respectively. RK3
840 requires three evaluations of the tendency G_c per time-step. RK3 is self-starting because
841 $\zeta^1 = 0$, while QAB2 must be started with a forward-backwards Euler step (the choice
842 $\chi = -1/2$ in (A3)). Equation (A2) is solved with a tridiagonal algorithm following a second-
843 order spatial discretization of $\partial_z \kappa_z^n \partial_z c^{m+1}$ — either once per time-step for QAB2, or three
844 times for each of the RK3’s three stages.

845 VerticallyImplicitTimeDiscretization permits longer time-steps when using fine vertical
846 spacing. Listing 14 illustrates the differences between vertically-implicit and explicit time
847 discretization using one-dimensional diffusion of by a top-hat diffusivity profile. The results
848 are shown in figure A1.

```
850
851     1 using Oceananigans
852     2
853     3 grid = RectilinearGrid(size=20, z=(-2, 2), topology=(Flat, Flat, Bounded))
854     4 time_discretization = VerticallyImplicitTimeDiscretization()
855     5 κ(z, t) = exp(-z^2)
856     6 closure = VerticalScalarDiffusivity(time_discretization; κ)
857     7 model = HydrostaticFreeSurfaceModel(; grid, closure, tracers=:c)
```

Listing 14: Diffusion of a tracer by a top hat tracer diffusivity profile using various time steps and time discretizations.

860 A2 The pressure correction method for momentum in NonhydrostaticModel

861 The NonhydrostaticModel uses a pressure correction method for the momentum equation
862 (6) that ensures $\nabla \cdot \mathbf{u} = 0$. We rewrite (6) as

$$\partial_t \mathbf{u} = -\nabla p + b \hat{\mathbf{z}} + \mathbf{G}_u + \partial_z (\nu_z \partial_z \mathbf{u}), \quad (\text{A5})$$

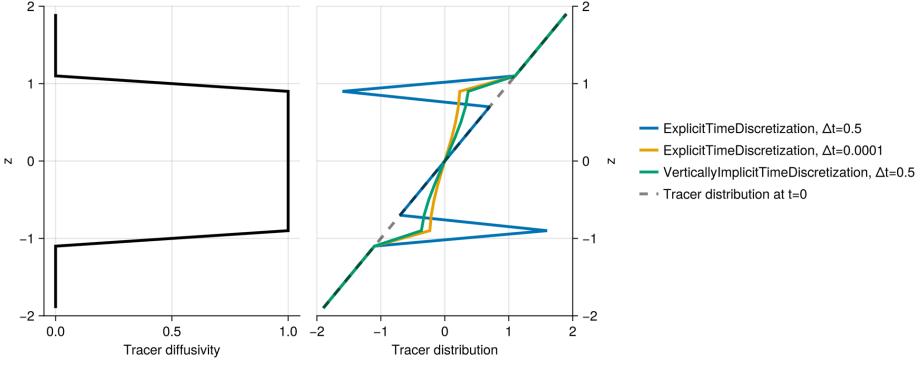


Figure A1: Simulations of tracer diffusion by a top hat diffusivity profile using various choices of time-discretization and time-step size. With a long time-step of $\Delta t = 0.5$, ExplicitTimeDiscretization is unstable while VerticallyImplicitTimeDiscretization is stable. Let the vertically-implicit solution depends on the long time-step $\Delta t = 0.5$, as revealed by comparison with ExplicitTimeDiscretization using $\Delta t = 10^{-4}$.

where, if specified, ν_z is the vertical component of the viscosity that will be treated with a vertically-implicit time discretization, ∇p is the total pressure gradient, and \mathbf{G}_u is the rest of the momentum tendency. We decompose p into a “hydrostatic anomaly” p' tied to the density anomaly ρ' , and a nonhydrostatic component \tilde{p} , such that

$$p = \tilde{p} + p', \quad \text{where} \quad \partial_z p' \stackrel{\text{def}}{=} b. \quad (\text{A6})$$

By computing p_h in (A6), we recast (A5) without b and with $\nabla p = \nabla p_n + \nabla_h p_h$. Next, integrating (A5) in time from t^m to t^{m+1} yields

$$\mathbf{u}^{m+1} = \mathbf{u}^m + \int_{t^m}^{t^{m+1}} [\mathbf{G}_u - \nabla \tilde{p} + \partial_z (\nu_z \partial_z \mathbf{u})] dt. \quad (\text{A7})$$

Next we introduce the predictor velocity $\tilde{\mathbf{u}}$, defined such that

$$(1 - \Delta t \partial_z \nu_z^m \partial_z) \tilde{\mathbf{u}} = \mathbf{u}^m + \int_{t^m}^{t^{m+1}} \mathbf{G}_u dt, \quad (\text{A8})$$

or in other words, defined as a velocity-like field that cannot feel nonhydrostatic pressure gradient $\nabla \tilde{p}$. Equation (A8) uses a semi-implicit treatment of vertical momentum diffusion which is similar but slightly different to the treatment of tracer diffusion in (A2),

$$\int_{t^m}^{t^{m+1}} \partial_z (\nu_z \partial_z \mathbf{u}) dt \approx \Delta t \partial_z (\nu_z^m \partial_z \tilde{\mathbf{u}}). \quad (\text{A9})$$

The integral in (A8) is evaluated with the same methods used for tracers — either (A3) for QAB2 or (A4) when using RK3. With a second-order discretization of vertical momentum diffusion, the predictor velocity in (A8) may be computed with a tridiagonal solver.

Introducing a fully-implicit time discretization for \tilde{p} ,

$$\int_{t^m}^{t^{m+1}} \nabla \tilde{p} dt \approx \Delta t \nabla \tilde{p}^{m+1}, \quad (\text{A10})$$

and inserting (A10) into (A8), we derive the pressure correction to the predictor velocity,

$$\mathbf{u}^{m+1} - \tilde{\mathbf{u}} = -\Delta t \nabla \tilde{p}^{m+1}. \quad (\text{A11})$$

The final ingredient needed to complete the pressure correction scheme is an equation for the nonhydrostatic pressure \tilde{p}_n^{m+1} . For this we form $\nabla \cdot (\text{A11})$ and use $\nabla \cdot \mathbf{u}^{m+1} = 0$ to obtain a Poisson equation for \tilde{p}_n^{m+1} ,

$$\nabla^2 \tilde{p}^{m+1} = \frac{\nabla \cdot \tilde{\mathbf{u}}}{\Delta t}. \quad (\text{A12})$$

Boundary conditions for equation (A12) may be derived by evaluating $\hat{\mathbf{n}} \cdot (\text{A7})$ on the boundary of the domain.

On RectilinearGrids, we solve (A12) using an eigenfunction expansion of the discrete second-order Poisson operator ∇^2 evaluated via the fast Fourier transform (FFT) in equispaced directions (Schumann & Sweet, 1988) plus a tridiagonal solve in variably-spaced directions. With the FFT-based solver, boundary conditions on \tilde{p}^{m+1} are accounted for by enforcing $\hat{\mathbf{n}} \cdot \tilde{\mathbf{u}} = \hat{\mathbf{n}} \cdot \mathbf{u}^{m+1}$ on boundary cells — which is additional and separate from the definition $\tilde{\mathbf{u}}$ in (A9). This alteration of $\tilde{\mathbf{u}}$ on the boundary implicitly contributes the appropriate terms that account for inhomogeneous boundary-normal pressure gradients $\hat{\mathbf{n}} \cdot \nabla \tilde{p}^{m+1} \neq 0$ to the right-hand-side of (A12) during the computation of $\nabla \cdot \tilde{\mathbf{u}}$.

A preconditioned conjugate gradient iteration may be used on non-rectilinear grids, including complex domains. For domains that immerse an irregular boundary within a RectilinearGrid, we have implemented an efficient, rapidly-converging preconditioner that leverages the FFT-based solver with masking applied to immersed cells. The FFT-based preconditioner for solving the Poisson equation in irregular domains will be described in a forthcoming paper.

A3 Time discretization of the HydrostaticFreeSurfaceModel

The HydrostaticFreeSurfaceModel uses a linear free surface formulation paired with a geopotential vertical coordinate that may be integrated in time using either a fully ExplicitFreeSurface, an ImplicitFreeSurface utilizing a two-dimensional elliptical solve, or a SplitExplicitFreeSurface. The latter free surface solver can also be used to solve the primitive equations with a non-linear free surface formulation and a free-surface following (z^*) vertical coordinate. For brevity, we describe here only the SplitExplicitFreeSurface, which is the most generally useful method. The SplitExplicitFreeSurface substeps the depth-integrated or “barotropic” horizontal velocity \mathbf{U}_h along with the free surface displacement η using a short time step while the depth-dependent, “baroclinic” velocities, along with tracers, are relatively stationary.

The barotropic horizontal transport \mathbf{U}_h is defined

$$\mathbf{U}_h \stackrel{\text{def}}{=} \int_{-H}^{\eta} \mathbf{u}_h \, dz, \quad (\text{A13})$$

where $\mathbf{u}_h = (u, v)$ is the total horizontal velocity and H is the depth of the fluid.

Similarly integrating the horizontal momentum equations (14) from $z = -H$ to $z = \eta$ yields an evolution equation for \mathbf{U}_h ,

$$\partial_t \mathbf{U}_h = -g(H + \eta) \nabla_h \eta + \int_{-H}^{\eta} \mathbf{G}_{uh} \, dz, \quad (\text{A14})$$

where \mathbf{G}_{uh} includes all the tendency terms that evolve “slowly” compared to the barotropic mode:

$$\mathbf{G}_{uh} = -(\mathbf{u} \cdot \nabla) \mathbf{u}_h - \mathbf{f} \times \mathbf{u} - \nabla \cdot \boldsymbol{\tau} + \mathbf{F}_h. \quad (\text{A15})$$

The evolution equation for the free surface is obtained by integrating the continuity equation (17) in z to obtain $\nabla \cdot \mathbf{U}_h = -w|_{z=\eta}$, and inserting this into (18) to find

$$\partial_t \eta = -\nabla_h \cdot \mathbf{U}_h. \quad (\text{A16})$$

The pair of equations (A14) and (A16) characterize the evolution of the barotropic mode, which involves faster time-scales than the baroclinic mode evolution described by equations (14). To resolve both modes, we use a split-explicit algorithm where the barotropic mode is advanced in time using a smaller time-step than the one used for three-dimensional baroclinic variables. In particular, a predictor three-dimensional velocity is evolved without accounting for the barotropic mode evolution, using the QAB2 scheme described by A3. We denote this “predictor” velocity, again, with a tilde as done in section A2.

$$(1 - \Delta t \partial_z \nu_z^m \partial_z) \tilde{\mathbf{u}}_h - \mathbf{u}_h^m \approx \int_{t^m}^{t^{m+1}} \mathbf{G}_{uh} dt. \quad (\text{A17})$$

We then compute the barotropic mode evolution by sub-stepping M times the barotropic equations using a forward-backward time-stepping scheme and a time-step $\Delta\tau = \Delta t/N$,

$$\eta^{n+1} - \eta^n = -\Delta\tau \nabla_h \cdot \mathbf{U}_h^n, \quad (\text{A18})$$

$$\mathbf{U}_h^{n+1} - \mathbf{U}_h^n = -\Delta\tau \left[g(H + \eta) \nabla_h \eta^{n+1} - \frac{1}{\Delta t} \int_{-H}^H \int_{t^m}^{t^{m+1}} \mathbf{G}_{uh} dt dz \right]. \quad (\text{A19})$$

The slow tendency terms are frozen in time during substepping. The barotropic quantities are averaged within the sub-stepping with

$$\bar{\mathbf{U}}_h = \sum_{n=1}^M a_n \mathbf{U}_h^n, \quad \bar{\eta} = \sum_{n=1}^M a_n \eta^n, \quad (\text{A20})$$

where M is the number of substeps per baroclinic step, and a_n are the weights are calculated from the provided averaging kernel. The default choice of averaging kernel is the minimal dispersion filters developed by Shchepetkin and McWilliams (2005). The number of substeps M is calculated to center the averaging kernel at t^{m+1} . As a result, the barotropic subcycling overshoots the baroclinic step, i.e. $M > N$ with a maximum of $M = 2N$. Finally, the barotropic mode is reconciled to the baroclinic mode with a correction step

$$\mathbf{u}_h^{m+1} = \tilde{\mathbf{u}}_h + \frac{1}{H + \eta} \left(\bar{\mathbf{U}}_h - \int_{-H}^H \tilde{\mathbf{u}}_h dz \right). \quad (\text{A21})$$

The barotropic variables are then reinitialized for evolution in the next barotropic mode evolution using the time-averaged $\bar{\eta}$ and $\bar{\mathbf{U}}_h$.

936 **Appendix B Table of numerical examples**

Description	Code	Visualization
2D turbulence using WENO(order=9) advection	listing 1	fig 1
2D turbulence with moving tracer source	listing 2	fig 1
DNS and LES of flow around a cylinder at various Re	listing 4	fig 2
DNS of cabbeling in freshwater	listing 5	fig 3
LES of the Eady problem with WENO(order=9)	listing 6	fig 4
Tidally-oscillating flow past Three Tree Point	listing 7	fig 5
Internal waves generated by tidal forcing over bathymetry	listing 8	fig 6
Comparison of vertical mixing parameterizations	listing 9	fig 7
Near-global ocean simulation with ClimaOcean	listing 10	fig 8
Visualization of the finite volume discretization	listing 11	fig 9
One-dimensional advection of a top-hat tracer profile	listing 13	fig 11
Tracer diffusion with various time discretizations	listing 14	fig A1

938 **References**

- 939 Arakawa, A. (1977). Computational design of the basic dynamical processes of the UCLA
940 general circulation model. *Methods in Computational Physics/Academic Press*.
- 941 Balsara, D., & Shu, C. (2000). Monotonicity preserving weighted essentially non-oscillatory
942 schemes with increasingly high order of accuracy. *Journal of Computational Physics*,
943 160(2), 405–452. doi: 10.1006/jcph.2000.6443
- 944 Besard, T., Foket, C., & De Sutter, B. (2018). Effective extensible programming: unleashing
945 Julia on GPUs. *IEEE Transactions on Parallel and Distributed Systems*, 30(4), 827–
946 841.
- 947 Bezanson, J., Edelman, A., Karpinski, S., & Shah, V. B. (2017). Julia: A fresh approach
948 to numerical computing. *SIAM review*, 59(1), 65–98.
- 949 Boccaletti, G., Ferrari, R., & Fox-Kemper, B. (2007). Mixed layer instabilities and restratification.
950 *Journal of Physical Oceanography*, 37(9), 2228–2250.
- 951 Bou-Zeid, E., Meneveau, C., & Parlange, M. (2005). A scale-dependent lagrangian dynamic
952 model for large eddy simulation of complex turbulent flows. *Physics of fluids*, 17(2).
- 953 Bryan, K. (1969). A numerical method for the study of the circulation of the world ocean.
954 *Journal of computational physics*, 135(2), 154–169.
- 955 Burns, K. J., Vasil, G. M., Oishi, J. S., Lecoanet, D., & Brown, B. P. (2020). Dedalus: A
956 flexible framework for numerical simulations with spectral methods. *Physical Review
957 Research*, 2(2), 023068.
- 958 Byrne, S., Wilcox, L. C., & Churavy, V. (2021). MPI.jl: Julia bindings for the Message
959 Passing Interface. In *Proceedings of the JuliaCon Conferences* (Vol. 1, p. 68). doi:
960 10.21105/jcon.00068
- 961 Callies, J., & Ferrari, R. (2018). Baroclinic instability in the presence of convection. *Journal
962 of Physical Oceanography*, 48(1), 45–60.
- 963 Chassignet, E. P., & Xu, X. (2017). Impact of horizontal resolution (1/12 to 1/50) on Gulf
964 Stream separation, penetration, and variability. *Journal of Physical Oceanography*,
965 47(8), 1999–2021.
- 966 Chassignet, E. P., & Xu, X. (2021). On the importance of high-resolution in large-scale
967 ocean models. *Advances in Atmospheric Sciences*, 38, 1621–1634.

- 968 Churavy, V. (2024). *KernelAbstractions.jl*. Retrieved from <https://doi.org/10.5281/zenodo.13773520>
- 969 Cox, M. D. (1984). *A primitive equation, 3-dimensional model of the ocean* (Tech. Rep. No. 1). Princeton, NJ: NOAA Geophysical Fluid Dynamics Laboratory.
- 970 Craik, A. D., & Leibovich, S. (1976). A rational model for Langmuir circulations. *Journal of Fluid Mechanics*, 73(3), 401–426.
- 971 Danilov, S., Sidorenko, D., Wang, Q., & Jung, T. (2017). The finite-volume sea ice–ocean model (fesom2). *Geoscientific Model Development*, 10(2), 765–789.
- 972 Dong, J., Fox-Kemper, B., Wenegrat, J. O., Bodner, A. S., Yu, X., Belcher, S., & Dong, C. (2024). Submesoscales are a significant turbulence source in global ocean surface boundary layer. *Nature Communications*, 15(1), 9566.
- 973 Eady, E. T. (1949). Long waves and cyclone waves. *Tellus*, 1(3), 33–52.
- 974 Forget, G., Campin, J.-M., Heimbach, P., Hill, C., Ponte, R., & Wunsch, C. (2015). Ecco version 4: An integrated framework for non-linear inverse modeling and global ocean state estimation. *Geoscientific Model Development*, 8(10), 3071–3104.
- 975 Ghosh, D., & Baeder, J. D. (2012). High-order accurate incompressible Navier–Stokes algorithm for vortex-ring interactions with solid wall. *AIAA journal*, 50(11), 2408–2422.
- 976 Godunov, S. K. (1959). A difference scheme for numerical solution of discontinuous solution of hydrodynamic equations. *Matematicheskii Sbornik*, 47, 271–306. (Translated by US Joint Publications Research Service, JPRS 7226, 1969)
- 977 Griffies, S. M., Adcroft, A., & Hallberg, R. W. (2020). A primer on the vertical lagrangian–remap method in ocean models based on finite volume generalized vertical coordinates. *Journal of Advances in Modeling Earth Systems*, 12(10), e2019MS001954.
- 978 Griffies, S. M., Pacanowski, R. C., & Hallberg, R. W. (2000). Spurious diapycnal mixing associated with advection in a z-coordinate ocean model. *Monthly Weather Review*, 128(3), 538–564.
- 979 Griffies, S. M., Stouffer, R. J., Adcroft, A. J., Bryan, K., Dixon, K. W., Hallberg, R., ... Rosati, A. (2015). A historical introduction to MOM. URL https://www.gfdl.noaa.gov/wp-content/uploads/2019/04/mom_history_2017.pdf, 9.
- 980 Häfner, D., Nuterman, R., & Jochum, M. (2021). Fast, cheap, and turbulent—global ocean modeling with GPU acceleration in Python. *Journal of Advances in Modeling Earth Systems*, 13(12), e2021MS002717.
- 981 Halliwell, G. R. (2004). Evaluation of vertical coordinate and vertical mixing algorithms in the HYbrid-Coordinate Ocean Model (HYCOM). *Ocean Modelling*, 7(3-4), 285–322.
- 982 Huang, N. E. (1979). On surface drift currents in the ocean. *Journal of Fluid Mechanics*, 91(1), 191–208.
- 983 Kärnä, T., Kramer, S. C., Mitchell, L., Ham, D. A., Piggott, M. D., & Baptista, A. M. (2018). Thetis coastal ocean model: discontinuous Galerkin discretization for the three-dimensional hydrostatic equations. *Geoscientific Model Development*, 11(11), 4359–4382.
- 984 Kiss, A. E., Hogg, A. M., Hannah, N., Boeira Dias, F., Brassington, G. B., Chamberlain, M. A., ... others (2020). Access-om2 v1. 0: a global ocean–sea ice model at three resolutions. *Geoscientific Model Development*, 13(2), 401–442.
- 985 Korn, P., Brüggemann, N., Jungclaus, J. H., Lorenz, S., Gutjahr, O., Haak, H., ... others (2022). Icon-o: The ocean component of the icon earth system model—global simulation characteristics and local telescoping capability. *Journal of Advances in Modeling Earth Systems*, 14(10), e2021MS002952.
- 986 Leclair, M., & Madec, G. (2011). z-coordinate, an arbitrary lagrangian–eulerian coordinate separating high and low frequency motions. *Ocean Modelling*, 37(3-4), 139–152.
- 987 Lilly, D. K. (1983). Stratified turbulence and the mesoscale variability of the atmosphere. *Journal of the Atmospheric Sciences*, 40(3), 749–761.
- 988 Marshall, J., Adcroft, A., Hill, C., Perelman, L., & Heisey, C. (1997). A finite-volume, incompressible Navier Stokes model for studies of the ocean on parallel computers. *Journal of Geophysical Research: Oceans*, 102(C3), 5753–5766.

- 1023 Marshall, J., Hill, C., Perelman, L., & Adcroft, A. (1997). Hydrostatic, quasi-hydrostatic,
 1024 and nonhydrostatic ocean modeling. *Journal of Geophysical Research: Oceans*,
 1025 *102*(C3), 5733–5752.
- 1026 McDougall, T. J., & Barker, P. M. (2011). Getting started with TEOS-10 and the Gibbs
 1027 Seawater (GSW) oceanographic toolbox. *Scor/iaupo WG*, *127*(532), 1–28.
- 1028 Molemaker, M. J., McWilliams, J. C., & Capet, X. (2010). Balanced and unbalanced routes
 1029 to dissipation in an equilibrated eady flow. *Journal of Fluid Mechanics*, *654*, 35–63.
- 1030 Pawlak, G., MacCready, P., Edwards, K., & McCabe, R. (2003). Observations on the
 1031 evolution of tidal vorticity at a stratified deep water headland. *Geophysical Research
 1032 Letters*, *30*(24).
- 1033 Pearson, B. (2018). Turbulence-induced anti-Stokes flow and the resulting limitations of
 1034 large-eddy simulation. *Journal of Physical Oceanography*, *48*(1), 117–122.
- 1035 Petersen, M. R., Jacobsen, D. W., Ringler, T. D., Hecht, M. W., & Maltrud, M. E. (2015).
 1036 Evaluation of the arbitrary Lagrangian–Eulerian vertical coordinate method in the
 1037 MPAS-Ocean model. *Ocean Modelling*, *86*, 93–113.
- 1038 Phillips, N. A. (1956). The general circulation of the atmosphere: A numerical experiment.
 1039 *Quarterly Journal of the Royal Meteorological Society*, *82*(352), 123–164.
- 1040 Pressel, K. G., Kaul, C. M., Schneider, T., Tan, Z., & Mishra, S. (2015). Large-eddy
 1041 simulation in an anelastic framework with closed water and entropy balances. *Journal
 1042 of Advances in Modeling Earth Systems*, *7*(3), 1425–1456.
- 1043 Pressel, K. G., Mishra, S., Schneider, T., Kaul, C. M., & Tan, Z. (2017). Numerics and
 1044 subgrid-scale modeling in large eddy simulations of stratocumulus clouds. *Journal of
 1045 advances in modeling earth systems*, *9*(2), 1342–1365.
- 1046 Ramadhan, A., Wagner, G., Hill, C., Campin, J.-M., Churavy, V., Besard, T., ... Marshall,
 1047 J. (2020). Oceananigans.jl: Fast and friendly geophysical fluid dynamics on GPUs.
 1048 *Journal of Open Source Software*, *5*(53).
- 1049 Ringler, T., Petersen, M., Higdon, R. L., Jacobsen, D., Jones, P. W., & Maltrud, M. (2013).
 1050 A multi-resolution approach to global ocean modeling. *Ocean Modelling*, *69*, 211–232.
- 1051 Roquet, F., Madec, G., Brodeau, L., & Nycander, J. (2015). Defining a simplified yet
 1052 “realistic” equation of state for seawater. *Journal of Physical Oceanography*, *45*(10),
 1053 2564–2579.
- 1054 Roquet, F., Madec, G., McDougall, T. J., & Barker, P. M. (2015). Accurate polynomial ex-
 1055 pressions for the density and specific volume of seawater using the TEOS-10 standard.
 1056 *Ocean Modelling*, *90*, 29–43.
- 1057 Rozema, W., Bae, H. J., Moin, P., & Verstappen, R. (2015). Minimum-dissipation models
 1058 for large-eddy simulation. *Physics of Fluids*, *27*(8).
- 1059 Schumann, U., & Sweet, R. A. (1988). Fast Fourier transforms for direct solution of
 1060 Poisson’s equation with staggered boundary conditions. *Journal of Computational
 1061 Physics*, *75*(1), 123–137.
- 1062 Shchepetkin, A. F., & McWilliams, J. C. (2005). The regional oceanic modeling system
 1063 (ROMS): a split-explicit, free-surface, topography-following-coordinate oceanic model.
 1064 *Ocean modelling*, *9*(4), 347–404.
- 1065 Shu, C. (1997). *Essentially non-oscillatory and weighted essentially non-oscillatory schemes
 1066 for hyperbolic conservation laws* (ICASE Report No. 97-65). Institute for Computer
 1067 Applications in Science and Engineering, NASA Langley Research Center.
- 1068 Shu, C.-W. (2009). High order weighted essentially nonoscillatory schemes for convection
 1069 dominated problems. *SIAM review*, *51*(1), 82–126.
- 1070 Silvestri, S., Wagner, G. L., Campin, J.-M., Constantinou, N. C., Hill, C. N., Souza, A., &
 1071 Ferrari, R. (2024). A new WENO-based momentum advection scheme for simulations
 1072 of ocean mesoscale turbulence. *Journal of Advances in Modeling Earth Systems*, *16*(7),
 1073 e2023MS004130.
- 1074 Silvestri, S., Wagner, G. L., Constantinou, N. C., Hill, C. N., Campin, J.-M., Souza, A. N., ...
 1075 Ferrari, R. (2024). A GPU-based ocean dynamical core for routine mesoscale-resolving
 1076 climate simulations. *Authorea Preprints*. doi: 10.22541/essoar.171708158.82342448/
 1077 v1

- 1078 Smagorinsky, J. (1963). General circulation experiments with the primitive equations: I.
 1079 The basic experiment. *Monthly weather review*, 91(3), 99–164.
- 1080 Stone, P. H. (1971). Baroclinic stability under non-hydrostatic conditions. *Journal of Fluid
 1081 Mechanics*, 45(4), 659–671.
- 1082 Tsujino, H., Urakawa, S., Nakano, H., Small, R. J., Kim, W. M., Yeager, S. G., ... others
 1083 (2018). Jra-55 based surface dataset for driving ocean–sea-ice models (jra55-do). *Ocean
 1084 Modelling*, 130, 79–139.
- 1085 Umlauf, L., & Burchard, H. (2005). Second-order turbulence closure models for geophysical
 1086 boundary layers. a review of recent work. *Continental Shelf Research*, 25(7-8), 795–
 1087 827.
- 1088 Vanneste, J., & Young, W. R. (2022). Stokes drift and its discontents. *Philosophical
 1089 Transactions of the Royal Society A*, 380(2225), 20210032.
- 1090 Vreugdenhil, C. A., & Taylor, J. R. (2018). Large-eddy simulations of stratified plane
 1091 Couette flow using the anisotropic minimum-dissipation model. *Physics of Fluids*,
 1092 30(8).
- 1093 Wagner, G. L., Chini, G. P., Ramadhan, A., Gallet, B., & Ferrari, R. (2021). Near-inertial
 1094 waves and turbulence driven by the growth of swell. *Journal of Physical Oceanography*,
 1095 51(5), 1337–1351.
- 1096 Wagner, G. L., Hillier, A., Constantinou, N. C., Silvestri, S., Souza, A. N., Burns, K., ... oth-
 1097 ers (2024). Formulation and calibration of CATKE, a one-equation parameterization
 1098 for microscale ocean mixing. *Authorea Preprints*. doi: 10.48550/arXiv.2306.13204
- 1099 Warner, S. J., & MacCready, P. (2014). The dynamics of pressure and form drag on
 1100 a sloping headland: Internal waves versus eddies. *Journal of Geophysical Research: Oceans*, 119(3), 1554–1571.